

Tracking Data and the Big Data Bowl

Every player, ten times a second: how to read the NFL’s tracking files, turn dots into football (separation, the pocket, motion, safety rotation), and what the Big Data Bowl has taught the league

A broadcast comes back from a commercial with a replay. The receiver runs a comeback, the ball arrives, and the screen freezes on a glowing ring around him: *separation 3.1 yards*. Later the same broadcast shows a safety’s path as a bright line sweeping from one hash to the middle of the field the instant the ball was snapped. Where do those numbers and lines come from, how much should you trust them, and could you make them yourself? Yes, with about a hundred lines of Python. This chapter shows how. It also builds, for real, the safety-rotation chart that Disguise, Rotation, and Split-Field Coverage showed you as a finished picture.

You’ll be able to:

- Read a tracking file: say what one row is, what `x`, `y`, `s`, `a`, `dis`, `o` and `dir` mean, and where the point (0, 0) is on the field.
- Flip every play so the offense moves left to right, and line thousands of plays up on the snap or on the throw.
- Draw and animate any play from tracking frames with `gridiron`, the same library that draws every diagram in this book.
- Compute separation, pocket area, motion at the snap, the shell at the snap and safety rotation from raw coordinates, and say what each number misses.
- Sketch three projects of your own: a man/zone classifier, a route classifier and a disguise measure.
- Recognise the Big Data Bowl’s themes since 2019, and what the winning entries had in common.

i You’ll need

- **Tendencies: PROE, Personnel, Motion, and Coverage: coverage charting** (people watching film and labelling each play’s coverage, the source of the coverage columns in `nflverse` data and in the Big Data Bowl files) and the **neutral situation** filter. This chapter replaces “what someone charted” with “where everyone actually was.”
- **Disguise, Rotation, and Split-Field Coverage: disguise** (showing one coverage and playing another), **post-snap rotation** (safeties changing the shell after the snap) and **late rotation** (moving only at the snap, so the quarterback cannot see it coming). The rotation chart in that chapter is rebuilt here, step by step.
- **Working with `nflverse` Data in Python:** `pandas`, and loading `nflverse` data with `nflreadpy`.

The data this chapter uses

Part 13 shows its code, so read the cells as part of the text. Two kinds of data appear.

Season summaries from Next Gen Stats, via `nflverse`. The NFL publishes some tracking-derived numbers for every season, such as each receiver’s average separation and each quarterback’s average time to throw. `nflverse` mirrors them, and `nfl().load_nextgen_stats()` loads them. These are real.

Frame-by-frame tracking, simulated. The raw files, with every player’s position ten times a second, are released only through the **Big Data Bowl** on Kaggle. You need a free account and must accept the competition’s rules to download them, and those rules limit how the data may be shared, so this book does not ship them.¹ Instead, every frame-level figure here uses plays simulated by `gridiron` and written out **in the exact format of the Big Data Bowl files** (same columns, same coordinate system, same quirks), labelled *simulated*, *BDB-format* wherever they appear. That matters more than it sounds: every function in this chapter runs unchanged on the real files. Wherever a figure is simulated, a cell marked “not run here” follows it with the code for the real data, which you can run once you have downloaded a Big Data Bowl release into `data/bdb/`.

Two kinds of code are kept out of your way. The **simulator** (the plumbing that fakes a Kaggle download) and the **drawing code** (font sizes, label positions) are folded away on the website under “Show the simulator” and “Show the plotting code”, and left out of the printed book. Every line that loads, reshapes or measures tracking data is shown.

¹Big Data Bowl data is distributed through Kaggle under each competition’s rules, which require accepting the rules before download and restrict use of the data; see the “Rules” and “Data” tabs of, for example, <https://www.kaggle.com/competitions/nfl-big-data-bowl-2025>. The exact terms differ by edition, so read them for the release you download.

```

import os
import tempfile
from pathlib import Path

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from matplotlib.lines import Line2D
from matplotlib.patches import FancyArrowPatch, Polygon
from IPython.display import Markdown

from gridiron import (Play, Field, offense, defense, routes as r, bdb, style, show_animation,
                      draw_frame, draw_marker, FIELD_WIDTH)
from gridiron.data import nfl

style.apply()
FPS = 10 # NFL tracking: 10 frames per second
SIM = Path(tempfile.mkdtemp(prefix="bdb_sim_")) # our stand-in for a Kaggle download folder

```

Where tracking data comes from

Every NFL player wears two or three small radio chips in his shoulder pads. Twenty to thirty receivers mounted around the stadium listen for them, and from the timing of the signals a computer works out where each chip is, about ten times a second. More than one chip per player gives you not just *where* he is but *which way his shoulders face*. The ball can carry a chip too. The league calls the whole system, and the group of analysts that turns its output into statistics, **Next Gen Stats**.²

Tracking data is the result: the recorded position of every player, many times a second, with speed, acceleration and body angles derived for each moment. **Next Gen Stats** (also called **NGS**) is both the NFL's tracking system and the league group that publishes stats from it.

Why chips, when basketball and baseball use cameras? The NBA had camera tracking in every arena from the 2013-14 season, and baseball's Statcast followed in 2015.³ Football is the hard case for cameras: twenty-two large bodies pile on top of each other at the line, and a camera cannot see a guard who is buried under two defensive tackles. A radio chip does not care what is on top of it. The NFL ran a trial in seventeen stadiums in 2014 and went league-wide the next season; the data has been the backbone of the league's analytics ever since.⁴

What tracking data does *not* record matters as much as what it does:

- **No hands, eyes or feet.** Each player is one point (roughly the middle of his shoulder pads). A tracking file cannot tell you that a corner jammed a receiver, that a guard's hands were outside the frame, or where the quarterback was looking.
- **No assignments.** It records where the free safety *went*, not where he was *supposed* to go. Every "coverage" label you will meet in a tracking file was added by people watching film (PFF, in the Big Data Bowl files) or by a model trained on such labels.
- **A little noise.** Positions are accurate to within about six inches,⁵ not to the millimetre, and a player standing still seems to drift slightly from frame to frame. That is why speeds near zero are unreliable, and why careful analysts smooth positions before computing speeds and accelerations from them (you will see why, with numbers, in a moment).

²NFL Football Operations, "NFL Next Gen Stats" ("2-3" RFID tags in each player's shoulder pads, 20-30 ultra-wide band receivers in each venue, data captured 10 times a second; tags also on officials, pylons, sticks, chains and in the ball; developed with Zebra Technologies): <https://operations.nfl.com/gameday/technology/nfl-next-gen-stats/>.

³NBA: "NBA expands partnership with STATS LLC to unveil player tracking technology for all 30 NBA teams" (Sept. 5, 2013; SportVU "in every NBA arena beginning with the 2013-14 season"): <https://pr.nba.com/nba-stats-llc-partnership>. MLB: "Statcast was installed in all 30 MLB ballparks in 2015 after a partial trial run in 2014": <https://www.mlb.com/glossary/statcast>.

⁴On the 2014 trial in 17 stadiums, Kevin Seifert, ESPN, "Inside Slant: NFL Next Gen Stats opens new doors" (Dec. 11, 2014): https://www.espn.com/blog/nflnation/post/_id/153569/inside-slant-nfl-next-gen-stats-opens-new-doors; on sensors in every stadium from the 2015 season, CBS Sacramento, "All NFL players will now wear motion trackers on the field" (Sept. 1, 2015): <https://www.cbsnews.com/sacramento/news/all-nfl-players-will-now-wear-motion-trackers-on-the-field>.

⁵Zebra Technologies' published figure, a margin of error of "less than six inches", as reported by CBS Sacramento (Sept. 1, 2015): <https://www.cbsnews.com/sacramento/news/all-nfl-players-will-now-wear-motion-trackers-on-the-field>; the NFL's own page says "within inches".

🔥 Common misconception

“Tracking data knows the coverage.” It knows where eleven defenders stood and moved. Whether that was Cover 3 or quarters is an *interpretation*, and two careful charting services can disagree about it. The man/zone and coverage columns in the Big Data Bowl files are charted labels sitting next to the tracking, not part of it. Much of the best Big Data Bowl work is about teaching a model to make that interpretation from the dots.

What is public, and when

Three layers of tracking data reach the public, at different speeds:

What	Detail	When you get it
Broadcast graphics	One number on one play (speed, separation, completion probability)	Live
NGS season and weekly summaries	Per-player averages: separation, cushion (how far off the receiver his defender lined up), time to throw, expected rushing yards and more; on nextgenstats.nfl.com and in nflverse	During the season, weekly
Raw frames (x, y, speed, angles for all 22 players)	Big Data Bowl releases only: a slice of past seasons	Once a year, for the competition

So the raw frames you can work with always lag the game you watched by a season or more. The 2026 Big Data Bowl release covers the 2023 season; the release with the richest pre-snap data covers the first nine weeks of 2022 (the 2025 Big Data Bowl). Everything in this chapter is written to work on both. The 2027 edition, launched in October 2026, is a different kind of release: sensor tracking from the Scouting Combine, paired with game tracking for the same players.⁶

One row, one player, one tenth of a second

A tracking file is a long table. Each row is **one player in one frame**. A **frame** is one snapshot of everyone on the field at one instant; at ten frames a second, a five-second play produces about fifty frames, so with 22 players plus the ball a single play is over a thousand rows, and one week of games is millions.

The field as a grid

Before the columns make sense, you need the grid they live on. It is the same grid every diagram in this book has been drawn on since the first chapter.

⁶Kaggle, “NFL Big Data Bowl 2027” (“Uncover non-obvious linkages between Combine sensor tracking and regular-season NFL game performance”; \$100,000 prize pool; entry deadline January 6, 2027): <https://www.kaggle.com/competitions/nfl-big-data-bowl-2027>; Yahoo Sports, “NFL wants to know which Combine measurements actually matter” (Oct. 7, 2026; 2023-2025 Combine data plus Next Gen Stats tracking for the same players): <https://sports.yahoo.com/articles/nfl-data-competition-focus-matters-123000120.html>.




Figure 1: The tracking coordinate system. x runs the length of the field, 0 to 120 yards, so the goal lines sit at $x = 10$ and $x = 110$ and midfield at $x = 60$. y runs across it, 0 to 53.3 yards; the hash marks sit at about $y = 23.6$ and $y = 29.8$. The point $(0, 0)$ is a corner of the field at the back of one end zone. The player shown stands at $x = 35$, $y = 20$: his own 25-yard line if his team is going to the right, nearer the $y = 0$ sideline. Angles (orientation and direction) are measured in degrees clockwise from the $+y$ direction, so 90 degrees points toward $x = 120$ and 270 degrees toward $x = 0$. After standardising (next section), every offense moves toward $+x$, which puts the high- y sideline on the offense's left and the low- y sideline on its right.

Four facts from that picture will save you hours:

- **The origin is a corner, not the ball.** Yard line numbers on the field ("the 25") are not x values. The 25-yard line nearest $x = 0$ is $x = 35$; the one nearest $x = 120$ is $x = 85$. Everything you compute about the line of scrimmage starts by finding the ball's x at the snap.
- **Once the offense faces $+x$, high y is its left.** Stand behind the quarterback looking toward $x = 120$: $y = 53.3$ is on your left hand, $y = 0$ on your right. So a receiver split wide to the offense's right has a *smaller* y than the ball, and "w yards to the offense's right", the number every diagram in this book is drawn in, is `ball_y - y`. Keep this one in your head; half the sign errors in tracking code come from it.
- **The field is 160 feet wide, which is $53\frac{1}{3}$ yards.** The files round it to 53.3. That gap of a few centimetres never matters for football, but it does matter for code: use one constant everywhere, or a play flipped twice will not land where it started. `gridiron` uses the exact $160/3$.
- **Angles are compass-like, but not math-like.** School trigonometry measures angles counter-clockwise from the x axis. Tracking angles run *clockwise* from the $+y$ direction. To turn an angle into a direction vector, use `dx = sin(angle)`, `dy = cos(angle)`, not the other way round. Get this wrong and every receiver appears to run sideways. For an offense heading toward $+x$, 90 degrees is straight downfield, 270 straight back toward its own end zone, 0 toward its left sideline and 180 toward its right.

The columns

Now the file itself. Here is the simulated play this chapter follows from start to finish, the **chapter play**: an 11-personnel 2x2 set (one back, one tight end, three receivers, two on each side), with the tight end Y flexed off the line into the right slot and the slot receiver H, on the left, sent in motion across the formation. The defense is a nickel that shows two-high safeties and rotates at the snap into a Cover 1 robber look: one safety spins to the deep middle, the other comes down into the middle of the field as a robber, and four defenders play man: the two corners on the outside receivers, the nickel on H (so he runs across the formation with the motion), and the Will linebacker on the tight end. The Mike sits in the middle watching the back. The quarterback throws a comeback to Z, the receiver wide to the right, on time: the ball leaves his hand as Z reaches the top of his stem (the straight, vertical part of his route), before he breaks back toward the sideline.

The corners are the one thing the library cannot do for us. Its man-coverage helper glues a defender to his receiver within about a second and a half, which is fine for a nickel on a slot but wrong for an off-coverage corner, who starts seven yards away and gives ground on purpose so nobody runs past him. So the play below steers both corners through timed waypoints with a small helper, `timed_path()`: backpedal, turn and run, then plant and drive on the ball.

```

def timed_path(p, key, keyframes):
    """Move defender `key` through (t, d, w) keyframes, t in seconds after the snap."""
    t_now, wait = 0.0, 0.0
    cur = np.array([p.players[key].d, p.players[key].w])
    for t, d, w in keyframes:
        dist = float(np.hypot(d - cur[0], w - cur[1]))
        if dist < 0.05:
            # standing still: wait, then go
            wait += t - t_now
        else:
            p.path(key, [(d, w)], absolute=True, speed=dist / (t - t_now), delay=wait)
            wait = 0.0
        t_now, cur = t, np.array([d, w])
    return p

def chapter_play(target="Z", t_throw=2.1, los=45, diagram=False):
    """The chapter play: 11 personnel 2x2 vs a nickel two-high shell that rotates to Cover 1 robber.
    diagram=True draws each corner's and Z's path as one line (same spots, no timing)."""
    off = offense("gun_2x2")
    dfn = defense("nickel", "two_high", off)
    spot = {"NB": (5.5, -8.3), "WILL": (5.0, 6.0), "MIKE": (4.5, 0.5), # nickel over H, Will over Y
            "LCB": (7.0, -15.8), "RCB": (7.0, 15.8), "FS": (13.0, -6.0), "SS": (13.0, 6.0)}
    dfn = [q.moved(d=spot[q.key][0], w=spot[q.key][1]) if q.key in spot else q for q in dfn]
    p = Play(off, dfn, los=los, to_go=10, title="The chapter play")
    # before the snap: the safeties bounce in place, then H motions across the formation
    p.motion("FS", to=(13.3, -5.8), speed=0.1)
    p.motion("SS", to=(12.8, 5.9), speed=0.07)
    p.motion("H", to=(-2.2, 3.5), speed=6.0) # still moving at the snap
    # protection: tackles kick back, the line gives ground late; the back checks the A/B gap
    for k, pts in {"LT": [(-1.6, -1.0), (-3.2, -0.5), (-4.6, 0.6)], "LG": [(-1.2, -0.3), (-2.4, 0.2)],
                  "C": [(-1.0, 0.0), (-2.0, 0.0)], "RG": [(-1.2, 0.3), (-2.4, -0.2)],
                  "RT": [(-1.6, 1.0), (-3.2, 0.5), (-4.4, -0.5)]}.items():
        p.block(k, pts=pts)
    p.block("RB", pts=[(0.4, 0.5)]) # up and inside the left guard
    for k, pts in {"WDE": [(-2.0, -1.8), (-4.8, -0.8), (-7.6, 2.2)],
                  "WDT": [(-1.6, 0.0), (-3.0, 0.3)], "SDT": [(-1.5, 0.2), (-3.0, -0.2)],
                  "SDE": [(-2.0, 1.8), (-4.6, 0.9), (-7.4, -1.8)]}.items():
        p.rush(k, pts=pts, speed=4.4)
    p.dropback(3.0) # a shotgun three-step drop
    p.route("X", r.go(24))
    if diagram:
        p.route("Z", r.path((14, 0), (11.5, 2.5), (10.7, 3.0)), speed=6.4)
    else:
        p.route("Z", r.comeback(14), speed=6.4) # 14 up, back 2.5 toward the sideline
        p.route("Z", [(-0.8, 0.5)], speed=1.5) # ... and keep working back to the ball
    p.route("Y", r.seam(18)); p.route("H", r.arrow(3, 6))
    p.man("NB", "H"); p.man("WILL", "Y") # the nickel travels with the motion
    # off-man corners: backpedal, turn and run, keep a cushion; the right one plants on the comeback
    corners = {"LCB": [(1.0, 9.3, -15.8), (2.0, 15.4, -16.1), (3.4, 25.0, -16.2)],
               "RCB": [(1.2, 10.0, 15.6), (2.0, 14.0, 16.0), (2.5, 16.2, 16.3),
                       (3.1, 12.4, 17.8), (3.6, 11.3, 18.5)]}
    for key, kf in corners.items():
        if diagram:
            p.path(key, [(d, w) for _, d, w in kf], absolute=True)
        else:

```

```

        timed_path(p, key, kf)
    p.drop("MIKE", (5.5, 0.0), speed=2.0)
    p.drop("FS", (16.5, 0.0), speed=6.5)
    p.drop("SS", (10.0, 2.0), speed=5.5)
    p.pass_to(target, t=t_throw)
    return p

chapter_play(diagram=True).draw(legend=True, width_in=6.2, window=(-9.5, 29), lateral=(-20, 21))
plt.show()
play = chapter_play()

```

sits in the middle, eyes on the back
 # spins to the deep middle at the snap
 # comes down as the robber

the timed version, for the tracking files

The chapter play




Figure 2: The chapter play as a playbook diagram (simulated). Offense: H's dashed line is his motion across the formation before the snap; X runs a go, Y a seam, Z a comeback (14 yards up, then back toward the sideline) and H a short arrow route, an angled route to the flat; the back steps up to block; the dashed brown arrow is the throw to Z. Defense: the dotted lines tie the two man defenders the library steers to their receivers (the nickel, \$, on H; the Will, W, on Y); the corners' lines show their paths (a backpedal, a turn to run with X and Z, and then, for the right corner, a plant and drive back down on the comeback); the free safety spins from just outside the left hash to the deep middle, the strong safety comes down to about 10 yards as the robber, the Mike drifts into the middle, and four rush.

A `gridiron` play can write itself out as tracking rows. A folded helper, `as_bdb_rows()`, does that and then makes the rows look like a real Big Data Bowl 2025 file: it adds the sensor wobble real chips have, adds which way each player is *facing* (the simulator moves dots but does not turn shoulders), turns the play around so that this offense is going toward the **left** end zone, as half of all real plays do, and `write_week()` writes the three files a real download would contain: `tracking_week_1.csv`, `plays.csv` and `players.csv`.

```
write_week(SIM, [(101, play, "left", {"playDescription": "(simulated) QB pass short right to Z",
                                     "pff_passCoverage": "Cover-1", "pff_manZone": "Man"})])
print(sorted(f.name for f in SIM.iterdir()))
```

```
['players.csv', 'plays.csv', 'tracking_week_1.csv']
```

Now load it exactly as you would load a real week, and look at one player's rows: a few from the start of the recording, then the snap and the frames just after it.

```
raw = pd.read_csv(SIM / "tracking_week_1.csv")
cb = raw[raw["displayName"] == "LCB"].sort_values("frameId")
f_snap = int(cb.loc[cb["frameType"] == "SNAP", "frameId"].iloc[0])
show_frames = [1, 2, 3, f_snap, f_snap + 1, f_snap + 2, f_snap + 7]
print("club:", cb["club"].iloc[0], "| playDirection:", cb["playDirection"].iloc[0])
cols = ["frameId", "frameType", "x", "y", "s", "a", "dis", "o", "dir", "event"]
md_table(cb[cb["frameId"].isin(show_frames)][cols].fillna(""))
```

```
club: DEF | playDirection: left
```

frameId	frame- Type	x	y	s	a	dis	o	dir	event
1	BEFORE_SNAP	67.99	10.88	0.21	0	0	87.53	184.31	
2	BEFORE_SNAP	67.99	10.86	0.21	0	0.02	91.24	184.31	
3	BEFORE_SNAP	68	10.86	0.14	0.64	0.01	86.16	165.54	
38	SNAP	68.03	10.84	0.1	0.34	0.03	90.99	89.22	ball_snap
39	AFTER_SNAP	67.76	10.87	0.78	6.73	0.27	89.97	275.79	
40	AFTER_SNAP	67.54	10.87	1.58	8.03	0.22	89.18	269.25	
45	AFTER_SNAP	66.41	10.89	2.19	0.7	0.22	92.49	273.19	

Read the columns from left to right:

Column	Meaning	Units
gameId, playId	Which game and which play. Together they identify a play everywhere in the release.	
nflId, displayName, jerseyNumber, club	Who the row is about. The ball has its own rows, with displayName "football" and no nflId.	
frameId	The frame number, counting from 1 at the start of the play's recording.	0.1 s per frame
frameType	BEFORE_SNAP, SNAP or AFTER_SNAP (Big Data Bowl 2025 files).	
time	In the raw file, the wall-clock timestamp of the frame. (<code>gridiron.bdb.prepare()</code> , used below, replaces it with seconds from the snap: negative before, positive after.)	

Column	Meaning	Units
<code>playDirection</code>	Which end zone the offense was heading for: right (toward $x = 120$) or left .	
<code>x, y</code>	Position on the grid above.	yards
<code>s</code>	Speed.	yards per second
<code>a</code>	Acceleration (its size, not its direction).	yards per second ²
<code>dis</code>	Distance travelled since the previous frame.	yards
<code>o</code>	Orientation: the way the player's shoulders face.	degrees
<code>dir</code>	Direction: the way he is moving.	degrees
<code>event</code>	What happened on that frame, when anything did: <code>line_set</code> , <code>man_in_motion</code> , <code>ball_snap</code> , <code>pass_forward</code> , <code>pass_arrived</code> , <code>pass_outcome_caught</code> , <code>tackle</code> and so on.	

Speeds are in yards per second, which few people have a feel for. Multiply by about **2.05** to get miles per hour: a receiver at full speed runs about 9 to 10 yards a second (18 to 21 mph), and the fastest ball carriers each season touch about 11 (around 22 mph, the numbers broadcasts love). A quarterback's dropback is 3 to 4 yards a second; a corner's backpedal a little slower.

Notice three things in the corner's rows (he is the left cornerback, and on this play the offense is going toward $x = 0$, so *downfield* for the offense means *smaller x*). First, in the early, pre-snap frames he is standing still, yet `s` is not exactly zero and `dir` jumps around: that is the sensor wobble. When a player is not moving, his direction is meaningless, so any feature built on `dir` should ignore frames with `s` below about **1 yard a second** (the same threshold the motion rule uses later). Second, by seven frames after the snap he is moving at about 2.2 yards a second with `dir` near 273 (toward $x = 0$, away from the line: he is backpedalling), while his `o` stays near 92 (he is facing the offense). Direction and orientation point almost opposite ways. That is the most important subtlety in the file, and it gets its own section next.

Third, look at `a` just after the snap: 8.2 yards per second², more than one g. No cornerback accelerates like that. It is an artefact of the simulator, which starts a player at full backpedal speed instantly, and of computing acceleration as the change in speed from one tenth of a second to the next. Real chips produce a milder version of the same problem, because differencing *amplifies* noise: a position wobble of a centimetre or two becomes a speed wobble of a few tenths of a yard per second (each difference is divided by 0.1 s), and differencing speed again to get acceleration amplifies it once more. The cure is to smooth the positions before you difference them:

```
def speed_from_xy(g, window=1):
    """Speed (yd/s) from positions, after a centred rolling mean over `window` frames."""
    g = g.sort_values("frameId")
    x = g["x"].rolling(window, center=True, min_periods=1).mean()
    y = g["y"].rolling(window, center=True, min_periods=1).mean()
    return np.hypot(x.diff(), y.diff()) * FPS

cb = cb.assign(s_raw=speed_from_xy(cb, 1), s_smooth=speed_from_xy(cb, 5))
still = cb[cb["frameType"] == "BEFORE_SNAP"].iloc[5:] # standing still before the snap
print(f"standing still, unsmoothed: speed {still['s_raw'].mean():.2f} yd/s on average, "
      f"acceleration {(still['s_raw'].diff().abs() * FPS).mean():.1f} yd/s2")
print(f"standing still, smoothed over 5 frames: speed {still['s_smooth'].mean():.2f} yd/s, "
      f"acceleration {(still['s_smooth'].diff().abs() * FPS).mean():.1f} yd/s2")
```

```
standing still, unsmoothed: speed 0.26 yd/s on average, acceleration 1.4 yd/s2
standing still, smoothed over 5 frames: speed 0.07 yd/s, acceleration 0.4 yd/s2
```

A five-frame rolling mean (half a second) is the bluntest tool that works; analysts who need sharper timing use a Savitzky-Golay filter, which fits a short polynomial through each window and keeps quick changes of direction better. Either way, the release's own `s`, `a` and `dir` columns are already derived numbers, so when you compute your own from `x` and `y`, smooth first and say how.

Orientation versus direction

Orientation vs direction is the difference between where a player *faces* (`o`) and where he *goes* (`dir`). Most of the time they agree: a receiver sprinting downfield faces downfield. But football is full of players moving one way while looking another, and those are exactly the players you want to understand.

To draw the play we need it in a friendlier shape, and `gridiron.bdb.prepare()` does that in one call. It does three jobs that the next two sections build by hand: it **flips the play** so the offense moves toward +x, it marks each row as offense, defense or ball, and it **rewrites time** as seconds from the snap. The cell then adds a few columns of our own: the position letters for the pictures, `d` and `w` (yards downfield of the ball and yards to the offense's right of it, the frame every diagram in this book is drawn in), and `frame_rel`, frames from the snap. Every later section uses this one table.

```
trk = bdb.load_tracking(str(SIM / "tracking_week_*.csv"))      # the same call works on real files
plays_csv = pd.read_csv(SIM / "plays.csv")
one = bdb.prepare(trk, plays_csv, game_id=2026091300, play_id=101) # flipped: offense moves to +x
one["label"] = one["displayName"].map({k: q.label for k, q in play.players.items()}).fillna("")
snap_ball = one[(one["frameType"] == "SNAP") & (one["side"] == "ball")].iloc[0]
LOS, BALL_Y = float(snap_ball["x"]), float(snap_ball["y"])
one["d"], one["w"] = one["x"] - LOS, BALL_Y - one["y"]          # high y is the offense's left
one["frame_rel"] = np.round(one["time"] * FPS).astype(int)      # frames from the snap
EVENT_T = one.dropna(subset=["event"]).drop_duplicates("event").set_index("event")["time"]
T_THROW, T_ARR = float(EVENT_T["pass_forward"]), float(EVENT_T["pass_arrived"])
print(f"line of scrimmage x = {LOS:.1f}, ball y = {BALL_Y:.1f}; "
      f"throw {T_THROW:.1f} s and catch {T_ARR:.1f} s after the snap")
```

```
line of scrimmage x = 45.0, ball y = 26.7; throw 2.1 s and catch 3.5 s after the snap
```




Figure 3: Orientation (solid black arrow, where the shoulders face) against direction (dashed green arrow, where the player is moving), 0.6 seconds after the snap of the chapter play (simulated, BDB-format; the quarterback is Q). The quarterback faces downfield while dropping straight back: the two angles are opposite. The free safety faces the quarterback while running sideways and deeper to the middle of the field, his rotation: the angles differ by more than 90 degrees. The defensive end on the offense's left, rushing, faces and moves the same way. A model that used only direction would think the quarterback was running away from the play and the safety was turning his back to it. (The simulated quarterback's chest points straight downfield; a real right-handed quarterback sets with his left shoulder toward the target, so his *o* sits well off 90 and swings as he works through his reads.)

Where does this matter? Everywhere a defender reads the play. A corner in off coverage gives ground with his shoulders square to the line of scrimmage; the instant he *turns* (his *o* swings toward his *dir*) he has committed, and a receiver who breaks at that moment wins. A zone defender keeps his shoulders open to the quarterback as he drops; a man defender turns his shoulders to run with his receiver. A pass rusher who has been turned sideways by a blocker is moving toward the quarterback but facing the sideline. Remember, though, that *o* is the chest, not the eyes: a man corner looking at his receiver's hips and a zone corner reading the quarterback can have exactly the same *o*. Several Big Data Bowl entries have built their best features from the gap between the two angles.

Standardising play direction

Half of all plays go left. If you average where the free safety lines up on a thousand plays without fixing that, half of him is at $x = 70$ and half at $x = 50$, and the average puts him on the ball. The fix is the first thing every tracking analysis does.

Standardised play direction means rotating every play that went toward the left end zone by 180 degrees, so that all plays move left to right. A rotation needs **both** coordinates flipped and both angles turned around:

```
def standardise(df):
    """Rotate plays with playDirection == 'left' by 180 degrees, so every offense moves toward +x."""
    df = df.copy()
    left = df["playDirection"] == "left"
    df.loc[left, "x"] = 120 - df.loc[left, "x"]
    df.loc[left, "y"] = FIELD_WIDTH - df.loc[left, "y"] # flip y too: a rotation, not a mirror
    for col in ("o", "dir"):
        df.loc[left, col] = (df.loc[left, col] + 180) % 360
    df.loc[left, "playDirection"] = "right"
    return df
```

```
mine = standardise(trk)
theirs = bdb.normalize_direction(trk) # gridiron's version of the same function
print("largest difference from gridiron's version:",
      float((mine[["x", "y"]] - theirs[["x", "y"]]).abs().max().max()), "yards")
```

largest difference from gridiron's version: 0.0 yards

Why flip y as well as x ? Picture yourself standing behind the offense. If you flip only x , the offense now faces the other way but the receiver who was on the quarterback's left is still at the same y , and from the quarterback's new point of view that is his *right*. You have made a mirror image, and every left-handed fact about the play (which hash, which side the tight end is on, the boundary) is now wrong. Flipping both is a half-turn of the whole field: nothing changes sides.




Figure 4: Standardising one play (simulated, BDB-format), shown at the snap with five players: the two outside receivers X and Z, the quarterback and the two safeties. (a) As recorded, the offense is heading for the left end zone ($x = 0$); X, the receiver on the quarterback's left, is on the low- y side. (b) Flipping only x turns the offense around but mirrors the play: X is now on the quarterback's right. (c) Flipping x and y rotates the play: the offense heads right and X is still on its left, now on the high- y side. (d) The same snap in the book's diagram frame, d yards downfield and w yards to the offense's right of the ball, drawn offense-up: this is the frame every play diagram in the book uses.

Panel (d) is worth a second look. Every diagram in this book is authored in two numbers: d , yards downfield of the line of scrimmage, and w , yards to the offense's right of the ball. After standardising, those are just $x - \text{LOS}$ and $\text{ball}_y - y$ at the snap (the minus sign is the "high y is the offense's left" fact from the grid). That is not a coincidence: the library was built so that a play a coach draws and a play the chips record are the same kind of object, and one drawing routine serves both.

! Why it exists

Every question worth asking of tracking data is a comparison: this safety against that one, this week against last, this team against the league. Comparisons need a common frame, and the raw files don't give you one: the same coverage looks like two opposite pictures depending on which end zone the offense faced. Standardising costs nothing and loses nothing (it is reversible), but forget it once and every average in the analysis is meaningless.

Lining plays up: the snap and the throw

The second alignment problem is time. Frame 1 is wherever recording started, which differs from play to play: one play has three seconds of frames before the snap, another twelve. So the next step is to re-number frames **relative to an event**: `frame_rel = frameId - frameId of the event`. Frame 0 is then the event itself, negative frames come before it, positive after.

The snap is the natural zero for most questions: pre-snap disguise, motion, how fast the rushers get off the ball. But many passing questions are really about **the throw**: how open was the receiver *when the ball left the hand*, how close was the pressure *then*. Quarterbacks throw anywhere from about 1.5 to over 4 seconds after the snap, so a chart aligned on the snap smears the moment you care about across thirty frames. Aligned on the throw, it lines up.

```
rng = np.random.default_rng(5)
choices = {"H": (1.4, 1.9), "Y": (1.9, 2.5), "Z": (2.0, 2.4), "X": (2.6, 3.2)} # throw times, s
rows = []
for pid in range(40):
    target = rng.choice(list(choices))
    p = chapter_play(target=target, t_throw=float(rng.uniform(*choices[target])))
    ball = p.to_tracking(fps=FPS, play_id=pid).query("displayName == 'football'").copy()
    f_snap = int(ball.loc[ball["event"] == "ball_snap", "frameId"].iloc[0])
    f_throw = int(ball.loc[ball["event"] == "pass_forward", "frameId"].iloc[0])
    ball["since_snap"] = ball["frameId"] - f_snap
    ball["since_throw"] = ball["frameId"] - f_throw
    ball["depth"] = ball["x"] - p.loc
    rows.append(ball.assign(target=target))
balls = pd.concat(rows)
```




Figure 5: Forty simulated dropbacks (BDB-format), each line one play: how far downfield the ball is. Left, aligned on the snap. In the first few frames every line drops below zero: that is the shotgun snap travelling five yards back to the quarterback, then his three-yard drop. The ball then sits in his hand and leaves on a throw anywhere from about 1.4 to 3.2 seconds after the snap, so the flights are smeared across the chart. Right, the same plays aligned on the pass_forward event: every flight starts at frame 0, so flights can be compared directly. Each line climbs until the catch and then levels off with the receiver; how high it climbs is how far downfield the target was, and how long it climbs is how long the ball was in the air.

In code, alignment is a merge: find each play's event frame once, join it back, subtract.

```
def align(frames, event="ball_snap"):
    """Add frame_rel: frames since `event` (negative = before it). Plays without the event drop out."""
    ev = (frames.loc[frames["event"] == event, ["gameId", "playId", "frameId"]
          .drop_duplicates(["gameId", "playId"]).rename(columns={"frameId": "f0"}))
```

```

out = frames.merge(ev, on=["gameId", "playId"])
out["frame_rel"] = out["frameId"] - out["f0"]
return out.drop(columns="f0")

```

```

at_throw = align(mine, "pass_forward")
print("frames before the throw:", -at_throw["frame_rel"].min(),
      "| after:", at_throw["frame_rel"].max())

```

```
frames before the throw: 58 | after: 21
```

Those 58 frames before the throw are the whole recording up to it: about 3.7 seconds of pre-snap frames (the recording starts when the safeties begin to bounce and H prepares to go in motion) plus 2.1 seconds from the snap to the throw. On real files, the pre-snap part varies from play to play, which is the whole reason to align.

Two practical warnings. **Events are not always where you expect.** The snap event name has varied (`ball_snap`, and `snap_direct` for direct snaps to a back), a few plays are missing an event, and the 2025 files added a `frameType` column that marks the snap frame directly; `gridiron.bdb.prepare()` checks all three. **The 2026 files split each play at the throw.** The Big Data Bowl 2026 asked entrants to predict where players would go *while the ball was in the air*, so its `input_` files hold the frames up to the throw and its `output_` files hold the positions afterwards, with snake_case column names (`game_id`, `play_id`, `nfl_id`, `frame_id`, `play_direction`) instead of camelCase.⁷ There is no throw event to align on there: the throw is the seam between the two files, so the last `input_` frame of a play is your frame 0. `gridiron.bdb.load_tracking()` renames the columns, but stitching the two halves back together is your job.

i Go deeper: getting the real files

Download. Create a Kaggle account, open the competition page (for example kaggle.com/competitions/nfl-big-data-bowl-2025) and accept the rules on the website; the download is refused until you do. Then either use the page's "Download all" button or, with Kaggle's command-line tool (`pip install kaggle`, plus an API token from your Kaggle account settings saved as `~/.kaggle/kaggle.json`), run `kaggle competitions download -c nfl-big-data-bowl-2025` and unzip into `data/bdb/`.

Size. Each tracking week holds 23 rows per frame, ten frames a second, for every play of every game that week, so a week runs to millions of rows. Load one week at a time, keep only the columns you need (`pd.read_csv(..., usecols=[...])`), and filter to your plays before anything else.

Quirks to check before you trust a result. Plays with no snap event; plays whose recording starts less than a second or two before the snap; the football's rows, which have no `nflId` (code that groups by `nflId` silently drops the ball); and, in the 2026 files, the input/output split above.

Drawing and animating a play

Everything is now in place to look at a play. The `one` table from the orientation section is exactly the input the drawing functions expect, whether the frames came from Kaggle or from a simulation: one row per player per frame, standardised, with `side`, `label` and `time` from the snap.

```

show_play(one, [(-1.5, "-1.5 s: motion"), (0.0, "snap"), (T_THROW, f"+{T_THROW:.1f} s: the throw"),
               (T_ARR, f"+{T_ARR:.1f} s: the catch")],
          window=(-10, 29), lateral=(BALL_Y - 21.5, BALL_Y + 19.5))

```

⁷Big Data Bowl 2025 tracking columns (camelCase: `gameId`, `playId`, `nflId`, `displayName`, `frameId`, `frameType`, `time`, `jerseyNumber`, `club`, `playDirection`, `x`, `y`, `s`, `a`, `dis`, `o`, `dir`, `event`) and 2026 columns (snake_case `input_*/output_*` files with `game_id`, `play_id`, `nfl_id`, `frame_id`, `play_direction`, ... `x`, `y`, `s`, `a`, `dir`, `o`, plus `num_frames_output`, `ball_land_x`, `ball_land_y`), as used in public solution code: <https://github.com/ZekeWeng/BigDataBowl2025> and <https://github.com/lneuendorf/nfl-big-data-bowl-2026>; competition overview at <https://operations.nfl.com/gameday/analytics/big-data-bowl/>.




Figure 6: The chapter play, drawn from its tracking file (simulated, BDB-format; faint lines are each player's path so far). -1.5 s: two-high safeties, and H in motion from the left slot with the nickel (\$) running across the formation with him, a sign of man coverage. Snap: H is still moving; nothing has rotated yet. Throw: the rotation is finished (the left safety, FS, is in the deep middle and the right safety, SS, has come down to about 10 yards as the robber), and Z is at the top of his stem with the right corner still over the top of him, two yards away. Catch: Z has broken back toward the sideline and worked back to the ball, and the corner has planted and closed to about a yard and a half, on Z's hip.

On the website that is a video you can scrub; in print it is four frames. Either way, notice what the picture shows that a coach's diagram cannot: *timing*. The robber is in place by about one second. The quarterback releases the ball as Z reaches the top of his stem, before the break, trusting him to win it, and the ball is in the air for well over a second.

With a real Big Data Bowl 2025 download, the same picture of any play takes the same three calls.

```
# Download from https://www.kaggle.com/competitions/nfl-big-data-bowl-2025/data (accept the
# rules first) and unzip into data/bdb/. Files: games.csv, plays.csv, players.csv,
# player_play.csv and tracking_week_1.csv ... tracking_week_9.csv.
trk = bdb.load_tracking("data/bdb/tracking_week_1.csv")
plays = pd.read_csv("data/bdb/plays.csv")
```

```

players = pd.read_csv("data/bdb/players.csv")[["nflId", "position"]]
# pick a play by its description, e.g. the first charted Cover-1 dropback of week 1
in_week = plays["gameId"].isin(trk["gameId"].unique())
pick = plays[(plays["pff_passCoverage"] == "Cover-1") & in_week].iloc[0]
real = bdb.prepare(trk, plays, game_id=pick.gameId, play_id=pick.playId) # labels = jersey numbers
real = real.merge(players, on="nflId", how="left")
real["label"] = real["position"].fillna("") # or keep the numbers
print(pick.playDescription)
show_animation(real, times=[-1.5, 0.0, 1.5, 3.0])

```

Expect a real play to look messier than the simulated one: players drift during the cadence, routes round off instead of breaking at sharp corners, and linemen bunch so tightly that their markers overlap. That mess is information, and every feature below is written to survive it.

Feature 1: separation

Separation is the distance from a receiver to the nearest defender. Next Gen Stats measures it at one moment, when the ball arrives (the catch or the incompletion), and averages it over a season.⁸ With frames you can measure it at *every* moment, which tells you much more than one number: when the receiver got open, and whether the defender closed the gap while the ball was in the air.

```

RECEIVERS = ["X", "H", "Y", "Z"]

def separation(frames, receivers=RECEIVERS):
    """Distance from each receiver to the nearest defender in every frame, and who that defender is."""
    dfn = frames.loc[frames["side"] == "def", ["frameId", "displayName", "x", "y"]]
    out = []
    for key in receivers:
        me = frames.loc[frames["displayName"] == key, ["frameId", "time", "x", "y"]]
        pair = me.merge(dfn, on="frameId", suffixes=("", "_def"))
        pair["sep"] = np.hypot(pair["x"] - pair["x_def"], pair["y"] - pair["y_def"])
        best = pair.loc[pair.groupby("frameId")["sep"].idxmin()]
        out.append(best.assign(receiver=key)[["frameId", "time", "receiver", "sep", "displayName"]])
    return pd.concat(out, ignore_index=True).rename(columns={"displayName": "nearest"})

sep = separation(one)
z = sep[sep["receiver"] == "Z"].set_index("time")["sep"]
mid = z[(z.index > T_THROW) & (z.index < T_ARR)]
at = sep[np.isclose(sep["time"], T_THROW) | np.isclose(sep["time"], T_ARR)]
md_table(at.pivot(index="receiver", columns="time", values="sep")
          .rename(columns={T_THROW: "at the throw", T_ARR: "at the catch"}).reset_index())

```

receiver	at the throw	at the catch
H	1.95	2.02
X	1.51	2.07
Y	1.44	1.42
Z	2	1.6

⁸NFL Next Gen Stats defines average separation as the distance in yards between the receiver and the nearest defender at the time of the catch or incompletion (nflverse Next Gen Stats data dictionary: https://nflreadr.nflverse.com/articles/dictionary_nextgen_stats.html).




Figure 7: Separation over time for each receiver on the chapter play (simulated, BDB-format): the distance to the nearest defender, from the snap to just after the catch. Most lines start high, because the corners and safeties begin 5 to 8 yards off. Z, the target, has about two yards when the ball is thrown: his corner is backpedalling over the top of him, keeping his cushion as off-man corners do. When Z breaks back toward the sideline the corner is still going the other way, and the gap jumps past four yards while the ball is in the air; the corner plants, drives, and is back to about a yard and a half by the catch. H’s line sits near 2 yards: the nickel who followed him across the formation trails him to the flat. The two vertical lines are the throw and the catch.

Two lessons hide in that chart.

Separation at the throw, at the catch and in between are three different numbers. The quarterback threw before the break, to a receiver who was not open yet: a timing throw, trusting Z to win it. Z was most open in mid-flight, about 4.3 yards, and the corner had cut that to about 1.6 by the catch. The NGS number you see on broadcasts is the last of the three. “Throwing receivers open” and “closing on the ball” are both real skills that a single snapshot hides, and measuring what happens while the ball is in the air was the whole point of the 2026 Big Data Bowl.

“Nearest defender” is not “covering defender.” On this play it is, because everyone near a receiver is in man. In zone coverage the nearest defender may be responsible for a different receiver entirely, and a receiver running past a linebacker who is dropping to his own zone looks covered for a frame or two. The cleaner measure, “distance to the defender assigned to him”, needs an assignment, and assignments are exactly what tracking does not record (the Big Data Bowl 2025 files add PFF’s charted matchups in `player_play.csv` for that reason).

What season-long separation really measures

NGS publishes each qualifying receiver’s average separation. Here is the 2025 season, plotted against how far downfield each player’s targets were.

```
ngs = nfl().load_nextgen_stats(seasons=list(range(2016, 2026)), stat_type="receiving").to_pandas()
season = ngs[(ngs["week"] == 0) & (ngs["season_type"] == "REG")]
rec25 = season[season["season"] == 2025].copy()
slope, icpt = np.polyfit(rec25["avg_intended_air_yards"], rec25["avg_separation"], 1)
r_air = np.corrcoef(rec25["avg_intended_air_yards"], rec25["avg_separation"])[0, 1]
by_player = season.pivot_table(index="player_gsis_id", columns="season", values="avg_separation")
yoy = by_player[[2024, 2025]].dropna() # receivers who qualified in both seasons
r_yoy = yoy.corr().iloc[0, 1]
rec25["resid"] = rec25["avg_separation"] - (icpt + slope * rec25["avg_intended_air_yards"])
```




Figure 8: Average separation at the catch or incompletion against average depth of target (intended air yards), 2025 regular season, every receiver who qualified in Next Gen Stats' season table. Real data: NFL Next Gen Stats via nflverse. Separation falls steadily the deeper the targets: deep routes are contested by design. Tight ends (orange) sit high and to the left because they are thrown short, often against linebackers. The labelled players are the season's least and most separated receivers and the Offensive Player of the Year, Jaxon Smith-Njigba.

The correlation between depth of target and separation is -0.69 : every extra yard of target depth costs about 0.12 yards of separation. So before you call a receiver "always open", ask how deep he runs. A fairer reading is the gap between a receiver's separation and the dashed line, separation *for his depth*. That is the same "over expected" idea you met with RYOE and CPOE, done with one variable instead of a model. Separation is also fairly stable: for receivers who qualified in both 2024 and 2025, the year-to-year correlation is 0.73, partly because route depth is a role that carries over.

Depth is not the only thing hiding in a season average. Separation reflects the **play-caller** as well as the receiver: motion, stacks, bunches, tight splits and play-action are all ways of manufacturing free releases and open throws, which is why receivers in motion-heavy offenses can post big numbers without being faster than anyone. And it reflects the receiver's **alignment**: a receiver who lines up outside and sees press coverage and the opponent's best corner every week will measure lower than a slot receiver running free against a linebacker.

```
th = rec25[rec25["player_display_name"] == "Tee Higgins"].iloc[0]
least = rec25.nsmallest(4, "avg_separation")[["player_display_name", "team_abbr", "avg_separation",
                                             "avg_intended_air_yards", "targets", "catch_percentage",
                                             "yards"]]
md_table(least.set_axis(["player", "team", "separation (yd)", "target depth (yd)", "targets",
                        "catch %", "yards"], axis=1), 2)
```

player	team	separation (yd)	target	depth (yd)	targets	catch %	yards
Tee Higgins	CIN	1.85		13.54	98	60.2	846
Mike Evans	TB	1.92		13.56	62	48.39	368
Christian Watson	GB	1.97		17.65	55	63.64	611

player	team	separation (yd)	target	depth (yd)	targets	catch %	yards
Alec Pierce	IND	2.07		19.09	84	55.95	1003

i Film room: Cincinnati Bengals 2025

The player: Tee Higgins, who had the lowest average separation of any qualifying receiver in 2025 (1.85 yards) and still caught 60% of his 98 targets for 846 yards, on targets averaging 13.5 yards downfield.⁹

What to notice: low separation is not the same as “not open.” Some receivers win *at the catch point*, with size, body position and hands, and their quarterbacks throw to them when they are covered. Watch a Bengals game and look at where the ball goes when Higgins is tightly covered on a fade or a back-shoulder throw. Then look at where he lines up: an outside receiver facing tight coverage most weeks is measured against a harder baseline than a slot receiver. The number describes how he wins, not whether he does; that is what makes a tracking stat useful, and also what makes it easy to misread.

Feature 2: the pocket

Pass Protection drew the pocket as the shape enclosed by the five linemen and the quarterback (its *convex hull*, the smallest shape with no dents that contains all six), and showed it breathing over a dropback. Here is that idea as a number you can compute on every frame of every play, with a second number that answers the question the quarterback is actually asking: how close is the nearest defender?

The hull itself is standard geometry, so it is folded away below as `hull()`. In your own environment, `scipy.spatial.ConvexHull(points).volume` gives the same area in one line (for two-dimensional points SciPy’s “volume” is the area); this book’s container doesn’t install SciPy, so the helper does it by hand, exactly as 04-02 did.

```
POCKET = ["LT", "LG", "C", "RG", "RT", "QB"]
rows = []
for fid, g in one[one["frame_rel"].between(0, 30)].groupby("frameId"):
    _, area = hull(g.loc[g["displayName"].isin(POCKET), ["x", "y"]].to_numpy())
    qb = g.loc[g["displayName"] == "QB", ["x", "y"]].to_numpy()[0]
    dfn = g.loc[g["side"] == "def", ["x", "y"]].to_numpy()
    rows.append((g["time"].iloc[0], area, np.hypot(*(dfn - qb).T).min()))
pocket = pd.DataFrame(rows, columns=["time", "area", "nearest"])
print(pocket.iloc[[0, 10, 20]].round(1).to_string(index=False))
```

```
time  area  nearest
0.0   13.7      6.0
1.0   22.1      6.0
2.0   18.3      3.4
```

⁹NFL Next Gen Stats season receiving table (2025 regular season, qualifying receivers), loaded with `nflreadpy.load_nextgen_stats(stat_type="receiving")` and cached in `data/cache/`; data retrieved October 2026. Higgins’s figures are this chapter’s calculation from that table.




Figure 9: The pocket on the chapter play (simulated, BDB-format). Left: the convex hull of the five linemen and the quarterback at the snap (blue, dotted), at +1.0 s (orange, dashed) and at the throw (green, solid and shaded), with the linemen, the back (R, stepped up inside the left guard) and the four rushers at the throw. Right: the hull’s area and the distance from the quarterback to the nearest defender, frame by frame. The hull grows as the tackles kick back and the quarterback drops, then shrinks again as the edge rushers push the tackles back toward him. The nearest defender stays about 6 yards away for the first one and a half seconds and is down to about 3 yards at the throw. Neither number alone is pressure: a big hull can still have a free rusher inside it.

Each of these numbers is a simplification, and it is worth saying of what. The hull measures how much space is *enclosed*, not *where* it is. Quarterbacks feel pressure from the front: a defensive tackle who drives the center two yards straight back takes away the spot the quarterback wanted to step up into, yet the hull barely changes, because the center then sits *inside* the shape the guards, tackles and quarterback already enclose (he stops being a corner of it, and the area loses only the sliver he used to add). The nearest defender ignores who is blocking him: a rusher two yards away and locked up by a tackle is no threat, a blitz three yards away and unblocked is. The tracking-based “win rates” ESPN publishes (pass rush win rate for one) solve that by asking a more specific question, whether a rusher has beaten his blocker within 2.5 seconds, and The Pass Rush shows how to compute get-off and time to pressure from the same frames.

Feature 3: motion at the snap

Motion and Shifts explained why offenses send a player across the formation, and why motion at the snap (a player still moving when the ball is snapped) has exploded: PFF counted shifts or motion on 63.9% of plays in 2025, up from 43.0% in 2018.¹⁰ Charting services decide by eye whether a play had motion, and their definitions differ: FTN’s `is_motion` flag in nflverse means motion “before or at” the snap, does not count shifts the way PFF’s “shift or motion” figure does, and runs lower (55% of 2025 plays against PFF’s 64%). Tracking data lets you write your own definition and apply it the same way to every play.

The rule is short. **A player is in motion at the snap if he is moving faster than some threshold in the snap frame.** The rulebook narrows who to check: only one player may be in motion at the snap, he must be an off-the-line (backfield) player, and he may not be moving toward the line. Linemen can’t be moving at the snap and the quarterback almost never is, so leave them out; their small movements are just noise. That rule also gives you a free data-quality check: if two or more skill players are above the threshold in the snap frame, you have found a shift that wasn’t set for a full second (a flag), a mislabelled snap frame, or a tracking glitch, and the play deserves a look.

A **shift** is different: several players reset, then everyone is still for a second before the snap. To find it you need the `line_set` event, the frame at which the *offensive line* has settled into its stances. Skill players are free to shift or go in motion after it (that

¹⁰Daire Carragher, PFF, “NFL pre-snap motion usage reaches new highs across the NFL” (Feb. 23, 2026): shift/motion rate 43.0% in 2018 and 63.9% in 2025: <https://www.pff.com/news/nfl-pre-snap-motion-usage-reaches-new-highs-across-the-nfl>. FTN’s `is_motion` definition (“motion occurred on the play before or at the time of the snap”): https://nflreadr.nflverse.com/articles/dictionary_ftn_charting.html.

is why the charted columns below are called `shiftSinceLineset` and `motionSinceLineset`), so a shift shows up as a player who has moved a few yards since `line_set` but is *stationary* at the snap.

```
SKILL = ["X", "H", "Y", "Z", "RB"]
THRESH = 1.0 # yards per second; wobble at rest is ~0.2-0.3
at_snap = one[(one["frameType"] == "SNAP") & one["displayName"].isin(SKILL)].set_index("displayName")
f_set = int(one.loc[one["event"] == "line_set", "frameId"].iloc[0])
start = one[(one["frameId"] == f_set) & one["displayName"].isin(SKILL)].set_index("displayName")
moved = np.hypot(at_snap["x"] - start["x"], at_snap["y"] - start["y"])
motion = pd.DataFrame({"speed at snap": at_snap["s"], "yards moved since line set": moved,
                      "in motion at snap": at_snap["s"] > THRESH,
                      "shifted": (at_snap["s"] <= THRESH) & (moved > 2.0)})
assert motion["in motion at snap"].sum() <= 1, "two players moving at the snap: check this play"
md_table(motion.reset_index().rename(columns={"displayName": "player"}))
```

player	speed at snap	yards moved since line set	in motion at snap	shifted
RB	0.2	0.03	False	False
X	0.1	0.05	False	False
H	5.89	12.49	True	False
Y	0.07	0.03	False	False
Z	0.1	0.04	False	False

The defense's answer to motion is just as measurable, and more interesting. If a defender moves across the field *with* the motion man, that is a man-coverage tell, the idea behind the man/zone indicator: in zone, defenders bump over a few yards or stay put; in man, someone travels. On the chapter play that someone is the nickel, the defender whose job is H.




Figure 10: Motion, measured (the chapter play; simulated, BDB-format). Top: speed of each skill player over the last three seconds before the snap and the first half-second after it. Everyone sits in the noise band near zero until the snap except H, who starts his motion about two seconds before the snap and is running at about 6 yards a second when the ball is snapped (frame 0): in motion at the snap. After the snap every receiver starts his route, so the grey lines climb too. Bottom: how far each defender has moved across the field since H started. The nickel moves almost exactly as far as H does (he travels with him: man coverage); the other defenders barely move until the snap.

That bottom panel is the seed of the first project in this chapter's last section: a man/zone classifier is, at heart, a careful version of "did anyone follow?". With the real files you can check your rule against the charters: the Big Data Bowl 2025 `player_play.csv` has `inMotionAtBallSnap`, `motionSinceLineset` and `shiftSinceLineset` for every player on every play.

```
pp = pd.read_csv("data/bdb/player_play.csv")
snap = trk[trk["frameType"] == "SNAP"].merge(players, on="nflId")
snap = snap[snap["position"].isin(["WR", "TE", "RB", "FB"])]
mine = snap.assign(my_motion=snap["s"] > 1.0)[["gameId", "playId", "nflId", "my_motion"]]
keys = ["gameId", "playId", "nflId"]
check = mine.merge(pp[keys + ["inMotionAtBallSnap"]], on=keys)
print(pd.crosstab(check["my_motion"], check["inMotionAtBallSnap"].fillna(False)))
# then tune the threshold (try 0.6 to 2.0 yd/s) and look at the plays where you disagree
```

Feature 4: safety rotation, step by step

This is the figure Disguise, Rotation, and Split-Field Coverage showed as a finished picture. Its job is to show *when* a defense's real coverage appears. Here it is built in four steps.

Step 1: find the safeties, and follow them on one play

The first decision is the least obvious: **who counts as a safety?** Not whoever the roster calls one. Teams play three-safety nickel packages, corners rotate to the deep middle on some calls, and a safety walked down into the box before the snap is not playing safety on that play. A robust definition uses the data itself: **the two deepest defenders one second before the snap**. Then measure, for each frame, how deep each one is (yards past the line of scrimmage) and how far he is sideways from the middle of the field.

```
def safety_tracks(frames, los, ball_y, lo=-30, hi=20):
    """The two deepest defenders 1 s before the snap: depth and distance from the ball, by frame."""
    f = frames.copy()
    f["frame_rel"] = np.round(f["time"] * FPS).astype(int)
    pre = f[(f["side"] == "def") & (f["frame_rel"] == -10)]
    ids = pre.nlargest(2, "x")["displayName"]
    s = f[f["displayName"].isin(ids) & f["frame_rel"].between(lo, hi)].copy()
    s["depth"] = s["x"] - los
    s["from_middle"] = (s["y"] - ball_y).abs()
    return s

saf1 = safety_tracks(one, LOS, BALL_Y)
print("safeties found:", sorted(saf1["displayName"].unique()))
```

```
safeties found: ['FS', 'SS']
```




Figure 11: Step 1: the two safeties on the chapter play alone (simulated, BDB-format), from 3 seconds before the snap to 2 seconds after. Top: depth past the line of scrimmage. Bottom: sideways distance from the middle of the field (here the ball sits in the middle). Before the snap both sit about 13 yards deep and about 6 yards off the middle, just outside the hashes: two-high. At the snap they split: the left safety (FS) goes deeper and runs to the middle (his sideways distance falls to zero), the right safety (SS) comes down to about 10 yards and pinches toward the middle: one-high, with a robber.

Step 2: many plays, one chart

One play is an anecdote. The chart in 06-06 stacks sixty: thirty on which a two-high shell stayed two-high, and thirty on which it rotated to one-high. The generator (folded below, and the same one that chapter used, with the same random seed so the figures match) varies where the safeties start, which side the rotation goes to, how fast they move and, crucially, *when* they start moving (τ , in seconds; negative means before the snap). Rotated plays send one safety to the deep middle and bring the other down to robber depth; robber depth itself varies with the call and the down and distance, anywhere from about 8 to 12 yards.

The loop applies `safety_tracks()` from step 1 to each play, then labels each safety's **role**: on rotated plays, the one who finishes deepest "spun" to the middle and the other "came down"; on the others both "stayed".

```

rng = np.random.default_rng(2026)
taus = [0.0] * 30 + list(rng.uniform(0.0, 0.25, 22)) + list(rng.uniform(-1.0, -0.5, 8))
tracks = []
for i, kind in enumerate(["stay2"] * 30 + ["rot1"] * 30):
    t = sim_dropback(kind, rng, tau=taus[i], play_id=i + 1)
    s = safety_tracks(t, los=35, ball_y=t["ball_y"].iloc[0])
    tracks.append(s.assign(kind=kind))
saf = pd.concat(tracks, ignore_index=True)
end = saf[saf["frame_rel"] == 20].set_index(["playId", "displayName"])[["depth"]]
deepest = end.groupby(level=0).idxmax().map(lambda t: t[1])
saf["role"] = np.where(saf["kind"] == "stay2", "stayed",
                      np.where(saf["displayName"] == saf["playId"].map(deepest), "spun", "dropped"))

```




Figure 12: Step 2, the book's safety-rotation figure (simulated, BDB-format tracking; not real plays): the two safeties on 60 plays that all showed two-high, from 30 frames (3 seconds) before the snap to 20 frames after. Thin lines are single plays; thick lines are medians. Top: depth. Bottom: sideways distance from the middle of the field. On 30 plays the shell stayed two-high (blue): both safeties drift back and a little wider. On the other 30 it rotated: one safety ran to the middle (orange: his sideways distance falls to about zero) while the other came down to robber depth (green: depth falls to about 9 or 10 yards, and he pinches in toward the middle). 22 of the 30 rotations were set to start at the snap (late rotation); the other 8 start 5 to 10 frames early, and their thin lines peel away before the snap: rotations the quarterback could have seen. Those 8 pull the thick orange and green lines into a gentle bend about 8 frames before the snap; the sharp split comes right after it.

Read it as the shape to look for, not as evidence: a simulation gives back only the timings it was given. A rotation is a **divergence**. Before it starts, rotated and honest plays are indistinguishable; within about ten frames they are in completely different places. Everything a defensive coordinator means by “rotate late” is about where that divergence begins.

Step 3: when did it rotate?

The chart shows timing by eye. To compare defenses you need it as a number for every play: the *rotation onset*, the frame at which the safety who finishes in the middle starts moving there. The rule below first decides whether the play rotated at all (one safety ends within 2.5 yards of the middle and the other ends shallower than 12 yards, deep enough to count a robber at the 10 to 12 yards the disguise chapter describes), then finds the first frame at which the middle safety is a yard closer to the middle than he was during the cadence.

```
def rotation_onset(s):
    """Frames from the snap at which the safety who ends in the middle starts moving there.
    None if the play never rotated to one-high (s: one play's output of safety_tracks)."""
    last = s[s["frame_rel"] == s["frame_rel"].max()].set_index("displayName")
    post = last["from_middle"].idxmin()
    other = last.drop(index=post)
    if last.loc[post, "from_middle"] > 2.5 or other["depth"].iloc[0] > 12.0:
        return None
    g = s[s["displayName"] == post]
    base = g.loc[g["frame_rel"].between(-30, -20), "from_middle"].median()
    return int(g.loc[g["from_middle"] < base - 1.0, "frame_rel"].min())

per_play = (saf.groupby("playId").apply(rotation_onset, include_groups=False)
            .rename("onset").reset_index())
per_play["onset"] = pd.to_numeric(per_play["onset"]) # None (no rotation) becomes NaN
onsets = per_play.dropna().astype({"onset": int})
n_late, n_early = (onsets["onset"] >= 0).sum(), (onsets["onset"] < 0).sum()
print(f"{len(onsets)} of 60 plays flagged as rotations; "
      f"{n_late} started at or after the snap, {n_early} before it")
```

30 of 60 plays flagged as rotations; 22 started at or after the snap, 8 before it




Figure 13: Step 3: rotation onset for every play the rule flagged as a rotation (simulated, BDB-format). Each dot is one play, placed at the frame at which its middle safety had moved a yard toward the middle. The shaded band marks the last half-second before the snap: a quarterback who sees the rotation start there can still change the protection or his first read, but rarely the play. The rule found all 30 rotations and none of the 30 honest plays. The 22 late rotations register 3 to 5 frames after the snap, and the 8 early ones 4 to 8 frames before it; in both groups the rule fires two or three frames after the safety really started, the time it takes him to cover the yard.

With rotation onset as a column, the measurement 06-06 sketched really is a few lines of pandas. On real files the `defense` column would be `defensiveTeam` from `plays.csv`, and `per_play` would hold every dropback on which the defense *showed* two-high; here all sixty simulated plays belong to one made-up defense:

```
profile = (per_play.assign(defense="SIM")                # real files: defensiveTeam
            .groupby("defense")
            .agg(two_high_looks=("playId", "size"),
                 rotate_share=("onset", lambda s: s.notna().mean()), # how often it lies
                 median_onset=("onset", "median")))                # how late
md_table(profile.reset_index())
```

defense	two_high_looks	rotate_share	median_onset
SIM	60	0.5	4

The first number is *how often* the shell lies; the second is *how late* it changes. A defense that scores high on both (it rotates often, and its median onset sits at or after the snap) is the hardest to read before the snap, which is the point of late rotation.

Step 4: the same thing on real files

```
# BDB 2025: 2022 season, Weeks 1-9, every frame from the huddle break to the end of the play.
plays = pd.read_csv("data/bdb/plays.csv")
dropbacks = plays[plays["pff_passCoverage"].notna()]
frames, onset_rows = [], []
for week in range(1, 10): # one week at a time keeps memory sane
    trk = bdb.load_tracking(f"data/bdb/tracking_week_{week}.csv")
    trk = trk.merge(dropbacks[["gameId", "playId"]], on=["gameId", "playId"])
    for (g, pl), _ in trk.groupby(["gameId", "playId"]):
        one = bdb.prepare(trk, plays, game_id=g, play_id=pl) # standardised, time 0 = snap
        snap = one[(one["frameType"] == "SNAP") & (one["side"] == "ball")]
        if snap.empty or one["time"].min() > -3.0: # need three seconds of pre-snap frames
            continue
        s = safety_tracks(one, float(snap["x"].iloc[0]), float(snap["y"].iloc[0]))
        frames.append(s.assign(gameId=g, playId=pl))
        onset_rows.append((g, pl, rotation_onset(s)))
saf = pd.concat(frames).merge(plays[["gameId", "playId", "defensiveTeam", "pff_passCoverage"]],
                             on=["gameId", "playId"])
per_play = pd.DataFrame(onset_rows, columns=["gameId", "playId", "onset"])
# Plot `saf` exactly as in step 2 (role from the coverage label and the deepest safety at +20);
# keep the plays that showed two-high at -1 s (shell_at, next section) and run the profile above
# grouped by defensiveTeam.
```

Expect the real version to be messier, in useful ways. Real safeties bounce, take false steps and react to motion, so the thin lines will be noisy. Plays start at different distances from the snap, so some will have fewer than thirty pre-snap frames. And the role labels will need the charted coverage, because a real defense can rotate and *stay* two-high: a quarters look that rolls into Cover 6 after the snap changes who has which deep zone without changing the shell. An “invert” Cover 2 is a sharper trap: the safeties come down to the flats while the corners bail to the deep halves, so after the snap the two deepest defenders are *corners*, and a feature that keeps following the two players who were deepest before the snap watches the wrong men. (The rotations Cover 3 Family taught, the buzz and the cloud, both end one-high, so this rule catches them.)

The shell at the snap, by depth and width

The chart above follows the shell through time. Often you want one label per play instead: **what shell did the defense show at the snap?** A tempting shortcut is a single depth: “if the second-deepest defender is deeper than about 11 yards, call it two-high.” It fails, because two-high safeties are not always deep. They usually sit 10 to 14 yards off the ball, and quarters safeties often align at 8 to 10 so they can fit the run; a one-high defense’s second-deepest man, a safety rolled down into the box or a robber, sits at 7 to 10. The depths overlap.

What separates the shells is **width**. In a two-high shell, the two deepest defenders sit on opposite sides of the ball, each on or outside the hash. In a one-high shell, the deepest defender sits in the middle. So the rule uses both:

```
def shell_at(defs, los, ball_y):
    """'two-high', 'one-high' or 'unclear' from one frame of defenders (columns x, y)."""
    d, w = defs["x"].to_numpy() - los, ball_y - defs["y"].to_numpy() # w: + = offense's right
    i1, i2 = np.argsort(d)[::-1][:2] # the two deepest
    if d[i2] >= 8 and w[i1] * w[i2] < 0 and min(abs(w[i1]), abs(w[i2])) >= 3:
        return "two-high" # both 8+ deep, opposite sides of the ball, on or outside the hashes
    if abs(w[i1]) <= 3:
        return "one-high" # the deepest man is in the middle
    return "unclear"
```

To see why width matters, here are 300 simulated snaps, half of each shell, drawn with realistic alignments (two-high safeties 8.5 to 14 yards deep, one-high down safeties 7 to 10.5):

```
rng = np.random.default_rng(22)
snaps = []
for i in range(300):
    two = i % 2 == 0
    if two:
        # two-high: one safety each side, quarters can be shallow
        sides = rng.choice([-1, 1]) * np.array([1, -1])
        d_s, w_s = rng.uniform(8.5, 14.0, 2), sides * rng.uniform(3.5, 8.5, 2)
    else:
        # one-high: a post safety and a rolled-down safety
        d_s = np.array([rng.uniform(12.0, 16.0), rng.uniform(7.0, 10.5)])
        w_s = np.array([rng.uniform(-2.5, 2.5), rng.choice([-1, 1]) * rng.uniform(3.0, 7.0)])
    d_cb = rng.uniform(5.0, 8.0, 2) # corners, 5 to 8 off, wide
    defs = pd.DataFrame({"x": 35 + np.r_[d_s, d_cb],
                        "y": FIELD_WIDTH / 2 - np.r_[w_s, [-15.0, 15.0]]})
    depths = np.sort(defs["x"].to_numpy() - 35)
    deepest = defs.loc[defs["x"].idxmax()]
    snaps.append({"truth": "two-high" if two else "one-high", "second_deepest": depths[-2],
                  "deepest_from_middle": abs(deepest["y"] - FIELD_WIDTH / 2),
                  "rule": shell_at(defs, 35, FIELD_WIDTH / 2)})
snaps = pd.DataFrame(snaps)
depth_only = np.where(snaps["second_deepest"] > 11, "two-high", "one-high")
print(f"depth cut at 11 yards: {(depth_only == snaps['truth']).mean():.0%} right; "
      f"depth and width: {(snaps['rule'] == snaps['truth']).mean():.0%} right")
```

depth cut at 11 yards: 64% right; depth and width: 100% right




Figure 14: Why the shell needs width as well as depth (300 simulated snaps, not real plays). Each dot is one snap: across, the depth of the second-deepest defender; up, how far the deepest defender is from the middle of the field. Depth alone (the dashed line) misreads every shallow two-high shell, such as quarters safeties at 9 or 10 yards, and every one-high look with a robber at 10 or more. Width separates them cleanly: in a two-high shell the deepest defender is a safety outside the hash, at least 3 yards from the middle; in a one-high shell he is the post safety, in the middle.

On these snaps the width rule is right every time, but only because the simulation was drawn from the same picture of football the rule encodes; the real files are the test. The rule is a first draft, and they will tell you where it needs work. Two refinements worth adding early: count only defenders *inside the widest receivers*, so that a corner bailing deep before the snap isn't mistaken for a safety, and classify the shell at several moments (two seconds before the snap, one second, the snap itself), because the change between them is the disguise.

Shells by team

The shell rule needs frames, which for a recent season you don't have. What nflverse does have is the *coverage played* on each dropback, charted after the snap. The chart below uses it for every defense in Weeks 1 to 9 of 2022, the slice of the 2022 season covered by the 2025 Big Data Bowl files: each team's share of two-high coverages (Cover 2, 2-Man, Cover 4 and Cover 6) among its charted dropbacks.¹¹

```
part = nfl().load_participation(seasons=[2022]).to_pandas()
pbp22 = (nfl().load_pbp([2022])
        .select(["game_id", "play_id", "defteam", "season_type", "week"]).to_pandas())
cov = part.merge(pbp22, left_on=["nflverse_game_id", "play_id", "right_on=["game_id", "play_id"]])
cov = cov[(cov["season_type"] == "REG") & (cov["week"] <= 9)
          & cov["defense_coverage_type"].notna() & (cov["defense_coverage_type"] != "PREVENT")]
cov["two_high"] = cov["defense_coverage_type"].isin(["COVER_2", "2_MAN", "COVER_4", "COVER_6"])
share = cov.groupby("defteam")["two_high"].mean().sort_values()
league = cov["two_high"].mean()
print(f"league: {league:.0%} two-high; lowest {share.index[0]} {share.iloc[0]:.0%}, "
      f"highest {share.index[-1]} {share.iloc[-1]:.0%}")
```

```
league: 41% two-high; lowest NE 23%, highest KC 60%
```

¹¹nflverse participation data (`load_participation`, 2022 season; coverage labels supplied by NGS for 2018-2022, FTN from 2023), joined to nflverse play-by-play for the defensive team and week; regular season, Weeks 1-9, dropbacks with a coverage label other than PREVENT (the 2022 labels have no separate Cover 9 or invert codes). Two-high = COVER_2, 2_MAN, COVER_4 and COVER_6. This chapter's calculation; data retrieved October 2026.




Figure 15: Each defense’s share of two-high coverages played (Cover 2, 2-Man, Cover 4, Cover 6) on charted dropbacks, 2022 regular season Weeks 1-9, the slice in the 2025 Big Data Bowl files. Real data: nflverse participation (coverage labels charted after the snap). The spread is wide, from New England and Miami, mostly one-high, to Kansas City and Buffalo, mostly two-high. These are coverages played, not shells shown: a defense that shows two-high and rotates late to Cover 3 counts as one-high here, so late-rotation teams would show more two-high at the snap than their dot suggests.

Read that chart with the chapter’s main lesson in mind: it is built from the coverage *played*, so it is a stand-in for the shell *shown*, and the two differ by exactly the disguise. Run `shell_at()` on the 2025 Big Data Bowl frames one second before each snap, put the share of two-high *looks* next to these dots, and the teams whose two numbers differ most are the ones that rotate. That comparison is Project 3. (The 2022 Seahawks in the chart predate Mike Macdonald, who coached Seattle from 2024.)

¹²NGS’s coverage classifier was built with the Amazon ML Solutions Lab and “identifies the defense coverage scheme based on the player tracking data” (trained on charted plays; launched for the 2022 season): AWS Machine Learning Blog, “Identifying defense coverage schemes in NFL’s Next Gen Stats”: <https://aws.amazon.com/blogs/machine-learning/identifying-defense-coverage-schemes-in-nfls-next-gen-stats/>.

¹³James Reber and the Next Gen Stats Analytics Team, NFL.com, “The coverage that turned the Seahawks into Super Bowl winners and is taking the NFL by storm” (2026): split-safety coverage on 32.9% of dropbacks in 2018 and 42.0% in 2025; Cover 4 at 17.2% in 2025; Seattle at 54.8% of pass plays from Week 14 through Super Bowl LX, “the highest rate of any playoff team in that span”, and 46% “in the nine games Love missed”: <https://www.nfl.com/news/the-coverage-that-turned-the-seahawks-into-super-bowl-winners-and-is-taking-the-nfl-by-storm>. The FTN-based figure (43.3% two-high among labelled 2025 regular-season dropbacks) is the course’s calculation from nflverse participation data.

The season: Mike Macdonald's Seahawks won Super Bowl LX in a season in which Next Gen Stats, whose coverage labels come from a model run on tracking data,¹² found the league in split-safety (two-high) coverage on 42.0% of dropbacks, the highest rate since its tracking of coverage began in 2018, with quarters (Cover 4) alone at 17.2%. Seattle went further. From Week 14 through the Super Bowl, the Seahawks played two-high on 54.8% of pass plays, the highest rate of any playoff team in that span, up from 46% in the nine games their safety Julian Love missed.¹³

What to notice: those numbers come from tracking; the coverage labels in nflverse for 2025 come from FTN's human charters. The two broadly agree for the league (FTN's labels give about 43% two-high for 2025), but they are different measurements of the same thing, and a team's number can differ between them. When a broadcast quotes a coverage rate, ask whose. Then watch Seattle's two safeties in the last second before the snap: how often does the shell you saw during the cadence survive the first second after it? The two-high to one-high case study shows what Macdonald asks them to do, and this chapter's shell rule and onset measure are how you would time it.

Three projects to try

The features above are building blocks. Here are three projects that put them together, each with the real labels you would train or test against, a sketch on simulated plays, and the trap that catches most first attempts.

Project 1: man or zone?

The question: from tracking alone, was the coverage man or zone? **The labels:** `pff_manZone` in the Big Data Bowl 2025 `plays.csv`; `team_coverage_man_zone` in the 2026 `supplementary_data.csv`. **The idea:** in man coverage a defender's job is a *person*, so he stays close to one receiver wherever that receiver goes. In zone his job is a *place*, so receivers run through his area and away from him. One feature captures much of that: for each defender, find the receiver nearest him at the snap, then measure how far apart they are a second or two later. Call it **stickiness**; small means man.

```
def stickiness(t, defenders=("LCB", "RCB", "NB", "WILL")):
    """Median distance (1.0-2.5 s after the snap) from each defender to the receiver nearest him
    at the snap, averaged over defenders. Small = defenders stayed with 'their' receivers."""
    snap = t[np.isclose(t["time"], 0.0)].set_index("displayName")
    later = t[t["time"].between(1.0, 2.5)]
    vals = []
    for d_ in defenders:
        gap = np.hypot(snap.loc[RECEIVERS, "x"] - snap.loc[d_, "x"],
                      snap.loc[RECEIVERS, "y"] - snap.loc[d_, "y"])
        mate = gap.idxmin()
        a_ = later[later["displayName"] == d_].set_index("time")
        b_ = later[later["displayName"] == mate].set_index("time")
        vals.append(np.hypot(a_["x"] - b_["x"], a_["y"] - b_["y"]).median())
    return float(np.mean(vals))

rng = np.random.default_rng(55)
sims, route_rows = [], []
for pid in range(240):
    kind = "man" if pid % 2 == 0 else "zone"
    t, ran = sim_coverage(kind, rng, pid)
    sims.append((pid, kind, stickiness(t)))
    ball_y = float(t.loc[(t["displayName"] == "football") & np.isclose(t["time"], 0.0), "y"].iloc[0])
    for key, name in ran.items():
        g = t[t["displayName"] == key][["time", "x", "y"]]
        route_rows.append(g.assign(playId=pid, receiver=key, route=name, ball_y=ball_y))
mz = pd.DataFrame(sims, columns=["playId", "label", "stickiness"])
cuts = np.linspace(1, 10, 181)
```

```
acc = [(mz["stickiness"] < c) == (mz["label"] == "man")).mean() for c in cuts]
best_cut, best_acc = cuts[int(np.argmax(acc))], max(acc)
print(f"best threshold {best_cut:.1f} yd -> {best_acc:.0%} of simulated plays classified correctly")
```

best threshold 2.0 yd -> 98% of simulated plays classified correctly




Figure 16: Project 1 on 240 simulated dropbacks (BDB-format), half Cover 1 man and half Cover 3 zone, with random routes. Stickiness is the average distance, one to two and a half seconds after the snap, between each of four underneath and outside defenders and the receiver nearest him at the snap. Man plays (blue) pile up below about 2 yards; zone plays (orange) spread out above, except where several receivers ran vertical routes: Cover 3 corners carry a vertical through their deep third, which looks exactly like man. The dashed line is the single best threshold.

The trap: on simulated plays, where a man defender follows his receiver perfectly, one number sorts 98% of plays correctly. Real football is much harder. Man corners play with leverage (lining up on one side of the receiver to take that side away) and give ground on purpose, as the chapter play's corners did; zone defenders in pattern-matching coverages *become* man defenders once routes declare; and some plays are man on one side and zone on the other. Expect a real-data version to need more features (orientation toward the receiver or toward the quarterback, how defenders react to motion, the safeties' movement) and a proper model. Above all, **split your training and test data by game, not by play**: plays from the same game share a defense, a game plan and a pair of quarterbacks, so a model tested on plays from games it trained on will look far better than it is.

Project 2: what route did he run?

The question: from a receiver's path alone, name his route. **The labels:** `routeRun` in the Big Data Bowl 2025 `player_play.csv`; in nflverse participation data, the targeted receiver's `route` (charted). **The idea:** standardise each path so that "toward his own sideline" is always positive, then describe it with a few numbers (how deep he got, where he ended up across the field, whether he came back toward the quarterback) and turn the route tree into rules. "His own sideline" is decided by which side of the *ball* he lined up on at the snap, not which half of the field he is in: on a play from the hash, a slot receiver can line up on the short side of the ball yet in the wide half of the field, and a rule based on the middle of the field would turn his out into an in.

```
def route_features(g):
    """Max depth, final depth and final sideways move (+ = toward his own sideline) of one path."""
    g = g.sort_values("time")
    d = g["x"].to_numpy() - g["x"].iloc[0]
    w = g["y"].to_numpy() - g["y"].iloc[0]
    out_sign = 1 if g["y"].iloc[0] < g["ball_y"].iloc[0] else -1 # right of the ball: out = -y
    out = -w * out_sign
    return pd.Series({"depth_max": d.max(), "depth_end": d[-1], "out_end": out[-1]})
```

```

def classify(f):
    """The route tree as if-then rules."""
    came_back = f["depth_max"] - f["depth_end"]
    if came_back > 0.6 and abs(f["out_end"]) < 4.5:      # turned back toward the quarterback
        return "comeback" if f["out_end"] > 1.0 else "curl/hitch"
    if f["depth_max"] >= 16:
        return "corner" if f["out_end"] > 4 else ("post" if f["out_end"] < -4 else "vertical")
    if f["out_end"] > 3:
        return "out"
    if f["out_end"] < -3:
        return "slant" if f["depth_max"] < 8 else "dig"
    return "other"

TRUE_CLASS = {"go": "vertical", "seam": "vertical", "post": "post", "corner": "corner", "out": "out",
              "dig": "dig", "slant": "slant", "curl": "curl/hitch", "hitch": "curl/hitch",
              "comeback": "comeback"}

paths = pd.concat(route_rows)
paths = paths[paths["time"] >= 0]
feats = paths.groupby(["playId", "receiver"]).apply(route_features, include_groups=False)
feats["predicted"] = feats.apply(classify, axis=1)
ran_route = paths.drop_duplicates(["playId", "receiver"]).set_index(["playId", "receiver"])["route"]
feats["truth"] = ran_route.map(TRUE_CLASS)
route_acc = (feats["predicted"] == feats["truth"]).mean()
n_other = (feats["predicted"] == "other").sum()
print(f"{len(feats)} simulated routes; rules agree with the true route on {route_acc:.0%}")

```

960 simulated routes; rules agree with the true route on 88%




Figure 17: Project 2: 960 simulated routes (BDB-format, each waypoint wobbled by about a yard), drawn from the receiver’s own starting point with his sideline to the right, grouped by the class the rules gave each one. Blue paths match the route the simulation ran; orange ones don’t. The biggest group of misses is ‘other’: curls and hitches that stopped without clearly coming back toward the quarterback, which the rules decline to name. The rest sit on a border between two classes: comebacks that drifted inside and read as curls, slants that ran deep enough to look like digs, outs that sank back like comebacks.

The rules get 88% of the simulated routes right, and the misses are instructive: they cluster at the borders between classes, where one yard of depth or one degree of angle decides the name. The largest group, the 95 routes the rules call “other”, are mostly curls and hitches whose wobbled path stopped without coming back the 0.6 yards the rule asks for. Lower the threshold and they become curls, but then some outs and digs that sag slightly at the end become curls too: every threshold trades one error for another. Real charters face the same borders and settle them with judgement (“he was supposed to run a curl”) that tracking cannot see.

The trap: treating the charted label as truth, and assuming it uses your names. Charting services have their own route vocabularies, and they change. nflverse’s 2022 participation data (charted by NGS) uses twelve names, **HITCH**, **FLAT**, **OUT**, **CROSS**, **GO**, **SCREEN**, **IN**, **SLANT**, **POST**, **ANGLE**, **CORNER** and **WHEEL**, while the 2025 data (charted by FTN) uses others, such as **QUICK OUT** and **DEEP OUT**, **IN/DIG**, **HITCH/CURL**, **SHALLOW CROSS/DRAW** and **SWING**.¹⁴ So a dig is an **IN** in one year and an **IN/DIG** in another, draws fall under a crossing label, and the flat, screen, angle, swing and wheel classes need rules of their own. Before you write a rule, run `value_counts()` on the label column of the release you have, map your classes onto it on paper, and check your model’s “errors” on film before you believe them.

Project 3: measuring disguise

The question: how often does a defense show one shell and play another, and how late does it change? **The labels:** `pff_passCoverage` in the Big Data Bowl 2025 `plays.csv` (the coverage played, charted after the snap). **The idea:** you have already built everything it needs. Label the shell two seconds before the snap and at the snap with `shell_at()`, compare both with the charted coverage, and add the rotation onset from step 3. Disguise, Rotation, and Split-Field Coverage shows what the result looks like on eight made-up defenses, and the code for the real files.

The trap: “two-high” is not one thing. A defense can show two-high and play Cover 6 or quarters honestly while its safeties creep and bounce; it can rotate *within* two-high; and charting labels themselves changed between seasons (nflverse participation data

¹⁴nflverse participation data, `route` column (the targeted receiver’s charted route), 2022 and 2025 seasons, loaded with `nflreadpy load_participation`; this chapter’s tabulation, data retrieved October 2026. The Big Data Bowl’s own `routeRun` labels come from PFF and may differ again: tabulate them before you use them.

switches from NGS to FTN coding in 2023, with different coverage names). Define your classes on paper before you write code, and keep the definition fixed across teams and seasons.

One more, if you want a fourth: expected yards after the catch. The Shanahan wide-zone case study compared nflverse’s expected YAC, which knows only where the ball was caught, with Next Gen Stats’ version, which knows where the defenders were. You can build the second kind yourself: take every completion’s catch frame, describe it with this chapter’s features (the receiver’s speed and direction, the distance and closing speed of the nearest two defenders, the separation at the catch), fit a model of the yards gained after the catch, and subtract. The trap is the same as in Project 1: split by game, and check that what your model calls “over expected” survives a second season.

Project 3 is, nearly word for word, the question the 2025 Big Data Bowl winners answered. Which brings us to the competition itself.

The Big Data Bowl

The **Big Data Bowl** is the NFL’s annual public analytics competition. Each year the league releases a slice of tracking data with a theme, and anyone (students, professionals, hobbyists) can enter. The finalists present at the NFL Scouting Combine, and the best ideas tend to show up a year or two later as Next Gen Stats metrics or broadcast graphics. It started in 2019; its co-creator is Michael Lopez, a former statistics professor who now leads the league’s football data and analytics group, and it was designed around a problem the league had: it was sitting on a uniquely rich data set that almost nobody outside the league and its clubs could study. Releasing slices to the public was a recruiting tool as much as a research one: the league says more than 75 Big Data Bowl participants have since been hired into data and analytics jobs in sports.¹⁵

Themes, 2019 to 2027

Edition	Theme	Data released	Winning idea, and what it led to
2019	Open questions on speed, routes and coverage	2017 season, first six weeks	Several entries modelled completion probability and “what if” throws to other receivers
2020	Rushing: predict the yards a run will gain from the hand-off	2017-2018 run plays; scored live on 2019 Weeks 13-17	The winning model became NGS’s expected rushing yards, the basis of RYOE
2021	Defending the pass	2018 season, every pass play	Winners classified man and zone coverage, then graded each defender’s effect on the play before and after the throw
2022	Special teams	2018-2020 punts, kickoffs and kicks	Winners computed a punt returner’s best path at every frame and graded his decisions against it
2023	Linemen on pass plays	2021 season, Weeks 1-8	Winners (“Between the Lines: How Do We Measure Pressure?”) measured pocket

¹⁵NFL Football Operations, “Big Data Bowl” (Mike Lopez, “senior director of football data & analytics at the NFL”; “more than 75 Big Data Bowl participants have been hired in data and analytics roles in sports”): <https://operations.nfl.com/gameday/analytics/big-data-bowl/>; Lopez as “co-creator of the Big Data Bowl”: AWS for M&E Blog, “The Big Data Bowl competition: Q&A with the NFL” (Mar. 1, 2023): <https://aws.amazon.com/blogs/media/the-big-data-bowl-competition-qa-with-the-nfl/>; his earlier post as a statistics professor at Skidmore College: The Bates Student, “Michael Lopez ‘04 on the NFL, Data, and Bates”: <https://thebatesstudent.com/22181/sports/michael-lopez-04-on-the-nfl-data-and-bates/>. Data slices for 2019-2024: 2019, all 91 games of 2017 Weeks 1-6 (Deshpande and Evans, arXiv:1910.12337); 2020, rushing plays scored on 2019 Weeks 13-17 (note below); 2021, all 2018 regular-season pass plays (<https://operations.nfl.com/updates/the-game/nfl-announces-third-annual-big-data-bowl/>); 2022, special teams 2018-2020 (<https://operations.nfl.com/updates/football-ops/nfl-announces-fourth-annual-big-data-bowl/>); 2023, linemen on pass plays, 2021 Weeks 1-8 (CMU workshop slides, <https://www.stat.cmu.edu/cmsac/conference/2022/workshop/slides.html>; Nguyen, Yurko and Matthews, below); 2024, tackling, 2022 Weeks 1-9 (<https://operations.nfl.com/updates/football-ops/nfl-announces-sixth-annual-big-data-bowl-competition-with-a-focus-on-tackling/>).

Edition	Theme	Data released	Winning idea, and what it led to
			pressure, and how long the pocket would last, at every frame
2024	Tackling	2022 season, Weeks 1-9	Winners (“Uncovering Missed Tackle Opportunities”) counted a missed tackle when a defender’s modelled tackle probability rose high and then fell with no tackle
2025	Pre-snap behaviour, and what it reveals	2022 season, Weeks 1-9, first edition with full pre-snap frames	“Exposing Coverage Tells in the Pre-Snap” won; NGS built a pre-snap disguise model
2026	Player movement while the ball is in the air	2023 season, Weeks 1-18; prediction track scored on late-2025 games	“Ghostbusters” (ghost defenders) won the analytics track
2027	Linking Combine sensor tracking to later regular-season performance	2023-2025 Combine sensor tracking (10 Hz) plus game tracking for the same players	Launched October 2026; entries close January 2027

The 2021-2024 winners are from contemporary reports of each edition.¹⁶ A few editions deserve a closer look, because their winning ideas changed how the league measures the game.

2020: from a competition to a broadcast stat. The rushing edition was a pure prediction contest: given the tracking frame at the handoff, predict how many yards the run would gain. The winners, Philipp Singer and Dmitry Gordeev, two Austrian data scientists competing as “The Zoo”, built a neural network that looked at every pair of an offensive player and a defender (how far apart they were, how fast they were closing), rather than at the players one at a time.¹⁷ Next Gen Stats adopted the approach for its expected rushing yards, which is why “rush yards over expected” now appears on broadcasts.

2025: reading the defense before the snap. The first edition with complete pre-snap frames asked what pre-snap movement reveals. The winners, NYU students Smit Bajaj and Vishakh Sandwar, built “Exposing Coverage Tells in the Pre-Snap”, a model of which pre-snap movements give a coverage away; Next Gen Stats then built a disguise model of its own, crediting their work.¹⁸ That is Project 3 above, and the tells catalogue in Watching Live is its human version.

2026: the ball in the air. The 2026 edition ran two competitions. In the prediction track, entrants were given every player’s tracking up to the throw and had to predict where each relevant player would be while the ball was in flight; the leaderboard was

¹⁶2021: Wei Peng, Marc Richards, Sam Walczak and Jack Werner, whose model classified “zone versus man-to-man coverage schemes” and assessed each defender before and after the throw (Associated Press via American Press, Mar. 19, 2021: <https://americanpress.com/?p=8915>). 2022: Robyn Ritchie, Brendan Kumagai, Ryker Moreau and Elijah Cavan, “quantifying the decision-making ability of the punt returner” by finding his optimal path (RStudio, “Tips for Getting Started With the NFL Big Data Bowl From the 2022 Winners,” Oct. 12, 2022: <https://www.r-bloggers.com/2022/10/tips-for-getting-started-with-the-nfl-big-data-bowl-from-the-2022-winners/>). 2023: Hassaan Inayatoli, Aaron White and Daniel Hovevar (University of Toronto), pocket pressure and “pocket longevity” (U of T Engineering News, Mar. 3, 2023: <https://news.engineering.utoronto.ca/?p=23871>). 2024: Matthew Chang, Katherine Dai, Daniel Jiang and Harvey Cheng, “Uncovering Missed Tackle Opportunities,” which flags a missed tackle when a defender’s tackle probability stays above 75% for at least half a second and then falls with no tackle made (Knowable Magazine, “How next-gen data analytics is changing American football,” 2024: <https://knowablemagazine.org/content/article/technology/2024/how-next-gen-data-analytics-changing-american-football>; title per TruMedia, “Matt Chang - NFL Big Data Bowl winner,” Mar. 12, 2024: <https://trumedianetworks.com/expected-value-podcast/2024/3/12/matt-chang-nfl-big-data-bowl-winner>).

¹⁷NFL Football Operations, “2020 Big Data Bowl results”: <https://operations.nfl.com/updates/the-game/2020-big-data-bowl-results>; NFL Next Gen Stats Analytics, “Next Gen Stats: Intro to Expected Rushing Yards” (2020): <https://www.nfl.com/news/next-gen-stats-intro-to-expected-rushing-yards>. The defense-by-offense pair design (“the first block of convolutions learns to work with defense-offense pairs of players”) is from the team’s write-up, “1st place solution The Zoo”, <https://www.kaggle.com/competitions/nfl-big-data-bowl-2020/writeups/the-zoo-1st-place-solution-the-zoo>, as quoted in a public reproduction: <https://github.com/juancamilocampos/nfl-big-data-bowl-2020>.

¹⁸NYU School of Professional Studies, “NYU duo crowned winners of the 2025 NFL Big Data Bowl”: <https://www.sps.nyu.edu/about/news-and-ideas/articles/preston-robert-tisch-institute-for-global-sport/2025/nyu-duo-crowned-winners-of-the-2025-nfl-big-data-bowl.html>; the NGS disguise model crediting the project: Next Gen Stats Analytics Team, “Next Gen Stats: New advanced metrics you need to know for the 2025 NFL season,” NFL.com (2025): <https://www.nfl.com/news/next-gen-stats-new-advanced-metrics-you-need-to-know-for-the-2025-nfl-season>.

scored live on games from Weeks 14 to 18 of the 2025 season. The analytics track was won by Lucca Ferraz of Rice University with “Ghostbusters: Back off man, I’m a Data Scientist!”, which built “ghost” defenders, a model of where an average defender would have gone, and graded real defenders against their ghosts.¹⁹ Notice the release’s quirk: the training files are the 2023 season, Weeks 1-18, even though some of the league’s announcement text mentions 2023 and 2024.

What the winners have in common

Read enough Big Data Bowl write-ups and a pattern shows up. The entries that win, and the ones the league later adopts, tend to share five habits.

1. **A football question first.** “Which pre-snap movements give away the coverage?” beats “here is a clustering of safety paths.” Judges are league and club football staff; a metric they can’t connect to a decision is a curiosity.
2. **An expectation, then a comparison.** Most winning metrics have the shape of RYOE: model what an average player would do in this exact situation, then measure the real player against it. Ghost defenders, a returner’s best path and a defender’s tackle probability are all the same idea.
3. **Relative features.** Positions relative to the ball, the line, the nearest opponent, and speeds relative to each other, not raw x and y. The 2020 winners’ pairwise design is the clearest case.
4. **Honest validation.** Does the metric predict something it was not built from? Is it stable from one half of the season to the other? Does it rank players in an order football people recognise, with surprises you can explain?
5. **Pictures.** Nearly every finalist animates plays. An animation of the metric changing frame by frame is how a coach decides whether to trust it, which is why this book’s library draws real and simulated plays the same way.

Go deeper

Earlier editions produced ideas you will meet again: a “hypothetical completion probability” for every receiver on a play, not only the one targeted, from the first edition (Sameer Deshpande and Katherine Evans);²⁰ and from the 2023 linemen edition’s data, STRAIN, a pass-rush measure built from how fast a rusher closes on the quarterback, by Quang Nguyen, Ronald Yurko and Gregory Matthews of Carnegie Mellon.²¹ Both were later published as academic papers, which are the best way to see what a strong entry looks like once it has been checked carefully.

¹⁹Kaggle, “NFL Big Data Bowl 2026 - Prediction” (<https://www.kaggle.com/competitions/nfl-big-data-bowl-2026-prediction>) and “NFL Big Data Bowl 2026 - Analytics” (<https://www.kaggle.com/competitions/nfl-big-data-bowl-2026-analytics>): theme, the two tracks, and the released training files `input_2023_w01..w18.csv` / `output_2023_w01..w18.csv`. Winner and project title per the NFL’s Big Data Bowl page, <https://operations.nfl.com/gameday/analytics/big-data-bowl>.

²⁰Sameer K. Deshpande and Katherine Evans, “Expected hypothetical completion probability,” *Journal of Quantitative Analysis in Sports* 16, no. 2 (2020): 85-94, <https://doi.org/10.1515/jqas-2019-0050>; the paper “elaborates on work done for the NFL 2019 Big Data Bowl contest” (arXiv:1910.12337).

²¹Quang Nguyen, Ronald Yurko and Gregory J. Matthews, “Here comes the STRAIN: Analyzing defensive pass rush in American football with player tracking data,” *The American Statistician* 78, no. 2 (2024): 199-208, <https://doi.org/10.1080/00031305.2023.2242442>. STRAIN uses the 2023 Big Data Bowl release (2021 Weeks 1-8); Carnegie Mellon, “Sports analytics: Quang Nguyen” (2025): <https://www.cmu.edu/dietrich/statistics-datascience/news/2025/sports-analytics-quang-nguyen.html>.

💡 Watch for it

Your first tracking chart. You can't download this week's tracking data, but you can practise the whole workflow on a play you watch live, and check part of it against real numbers within days.

- **Pick one play** during the game, ideally a dropback with motion. Write down the quarter, the clock, the down and distance, and three facts: the shell before the snap (two deep defenders on either side of the ball, outside the hashes, is two-high; one deep defender in the middle is one-high), whether anyone was moving at the snap, and where the ball went.
- **Watch one safety** from the moment the offense sets until the throw. Did he move before the snap or at it? Say "now" out loud when he starts to rotate; you are estimating rotation onset by ear.
- **After the game**, find the play in nflverse play-by-play (`load_pbp()` for the current season, matched on quarter and clock) to get the official description and the target.
- **Draw it** in `gridiron` the way you remember it: formation, motion, the routes you saw, the safety rotation. Export it with `to_tracking()` and run this chapter's `separation()`, `shell_at()` and `safety_tracks()` on it.
- **Check what you can.** Next Gen Stats' weekly receiving numbers (`load_nextgen_stats(stat_type= "receiving")`, weekly rows) include the targeted receiver's average separation and cushion for that game. Was your play typical for him?
- **Keep the file.** When that season's tracking appears in a future Big Data Bowl, run the same code on the real frames and see how good your eyes were.

Predict the play

The pictures use the course's position letters: on offense X, Z and H are wide receivers, Y the tight end, R the back and Q the quarterback; on defense E is an end, T a tackle, M and W the middle and weak-side linebackers, C a cornerback, \$ the nickel back, and FS and SS the safeties.

💡 Predict the play

Motion, and nobody follows. The offense is in the same 2x2 set as the chapter play and sends H across the formation from the left slot. The two frames below are 2 seconds before the snap and the snap itself, with the defenders' trails drawn thick so you can see who moved, and how far. Before the snap: is the defense more likely in man or zone? What would this chapter's motion and travel measures say? And with two safeties still at 13 yards at the snap, what do you still *not* know?




Figure 18: Drill 1 (simulated, BDB-format): H in motion from the left slot to the right, 2 seconds before the snap and at the snap. The thick orange lines are the defenders' trails; watch them, not the motion man.

Answer 1 is at the end of the chapter. Try it first.

💡 Predict the play

Read the chart. Here are the two deepest defenders on one simulated play, A and B, measured exactly as in step 1 of the rotation section. Before the snap, what shell did the defense show? What did it play after the snap? Did the change start early enough for the quarterback to use it, and what would `rotation_onset()` return?




Figure 19: Drill 2: the two deepest defenders on one play (simulated, BDB-format), 3 seconds before to 2 seconds after the snap. Top: depth. Bottom: sideways distance from the middle of the field. The shaded half is before the snap.

Answer 2 is at the end of the chapter. Try it first.

💡 Predict the play

Read the rows. One frame from the chapter play, 0.7 seconds after the snap, with three players marked A, B and C. Their rows from the standardised file are in the table under the picture. Which one is backedpedalling, which is running a route, and which is standing still? For which of them should you ignore `dir`?




Figure 20: Drill 3 (the chapter play, simulated, BDB-format), 0.7 seconds after the snap. Three players are marked A, B and C (green rings); everyone else is unlabelled. The offense is moving up the page.

player	x	y	s	o	dir
A	53.6	42.4	2.2	272.5	93.2
B	45.2	18.5	7.5	149	153.9
C	40.4	27.6	0.1	83.8	138.3

Answer 3 is at the end of the chapter. Try it first.

Takeaways

- **A tracking file is one row per player per frame**, ten frames a second, on a 120×53.3 -yard grid whose origin is a corner of the field. `s` and `a` are in yards per second ($\times 2.05$ for mph) and yards per second², `o` is where a player faces and `dir` where he moves, both in degrees clockwise from $+y$. Smooth positions before you compute your own speeds or accelerations.
- **Standardise and align before anything else**: rotate left-going plays by 180 degrees (flip `x` and `y`, turn both angles; high `y` is then the offense's left), then number frames from the event your question is about, the snap or the throw.
- **The same library draws real and simulated plays**. `bdb.prepare()` turns a raw file into frames; `show_animation()` and `frame_strip()` draw them; this book's `d` and `w` are just standardised `x` and `y` measured from the ball.
- **Features are simple geometry with hidden choices**: separation is "nearest defender" (not the covering one) and depends on when you measure it; pocket area ignores who is blocked and where the space is; motion needs a speed threshold; "safeties" are best defined as the two deepest defenders before the snap, and the shell needs their width as well as their depth. Season-long separation is mostly a route-depth stat, with the play-caller's fingerprints on it.
- **Rotation is a divergence in time**, and rotation onset turns "late rotation" into a number per play. Disguise is two numbers per defense: how often the shell lies, and how late it changes.

- **The Big Data Bowl** has released a new slice of tracking data with a new theme every year since 2019. Its best ideas (expected rushing yards, pocket pressure, missed-tackle opportunities, pre-snap disguise models, ghost defenders) share a shape: a football question, an expectation, a comparison, honest validation and a picture.

Connections

Builds on: Tendencies: PROE, Personnel, Motion, and Coverage (the charting labels that tracking is compared with) and Disguise, Rotation, and Split-Field Coverage (what the safety-rotation chart means). The pocket hull comes from Pass Protection, get-off and time to pressure from The Pass Rush, the comeback and its stem from The Quick Game, and RYOE from Expected Points, EPA, and Success Rate.

Sets up:

- Fangio’s Two-High Revolution and Macdonald’s Answer: the defenses this chapter’s rotation chart was built to read.
- Watching Live: the full catalogue of pre-snap tells, the human version of the 2025 Big Data Bowl.
- Film Study and Charting: charting your own games, and scoring your predictions against what happened.
- Scheme and Fantasy: separation and target depth as usage signals.

If you want to go deeper: the Big Data Bowl pages on Kaggle (each edition’s data description and the winning notebooks are public); Next Gen Stats’ own site, nextgenstats.nfl.com, for the published metrics and their definitions; Michael Lopez’s introduction to the *Journal of Quantitative Analysis in Sports* special issue on tracking data, a short, readable map of the field;²² and the academic versions of winning entries cited above. For working code, the public Big Data Bowl notebooks on Kaggle are the fastest way to see how other people standardise, smooth and feature-engineer the same files.

Answers to Predict the play

Answer 1.

Zone is the better bet. Nobody travelled with H: the Will linebacker who was over him slid two or three yards toward the middle and the Mike slid over, but no defender crossed the formation with him. In this chapter’s terms, every defender’s “yards moved across since motion began” stays small while H’s grows past 12, so the travel feature says *no man defender followed*. The two linebackers bumping over together is what zone defenses do to keep their areas balanced to the new strength.

What you still don’t know is the shell after the snap. Two safeties at 13 yards at the snap fit Cover 2, quarters or Cover 6, but they also fit a late rotation to Cover 3. Motion answered “man or zone?”; only the first frames after the snap answer “which zone?”. That is exactly why defenses that rotate late pair it with zone-looking motion responses: the motion tell and the shell tell can be made to point different ways.

Answer 2.

One-high shown, two-high played. Two seconds before the snap one defender sits deep in the middle (about 13 yards deep, within a yard or so of the middle) and the other is down at about 8 yards, several yards off to one side: a one-high shell. Two seconds after the snap both are 16 to 18 yards deep and about 10 yards either side of the middle: two deep halves, so Cover 2 (or 2-Man, with man underneath). It is not quarters: quarters safeties stay shallower, around 10 to 13 yards, and inside the slot receivers, reading them, rather than bailing to 15-plus yards near the numbers. This is a “spin” the other way from the chapter play, one-high to two-high.

The movement starts about **eight frames before the snap**, so a quarterback watching the safeties could see the shell changing during his cadence: not a late rotation. And `rotation_onset()` returns None, because it was written to find rotations *to* one-high (one safety finishing in the middle). That is a useful reminder about home-made features: they answer exactly the question you coded, so a play that moves the other way slips through silently. Extending it to spins toward two-high is a good first edit.

Answer 3.

A is backpedalling. He faces the backfield (o near 270 degrees, toward x = 0) while moving downfield (dir near 90) at about 2.2 yards a second: the left corner, giving ground in off coverage with a cushion of about four yards on X, before he turns to run with him.

²²Michael J. Lopez, “Bigger data, better questions, and a return to fourth down behavior: an introduction to a special issue on tracking data in the National Football League,” *Journal of Quantitative Analysis in Sports* 16, no. 2 (2020): 73-79, <https://doi.org/10.1515/jqas-2020-0057>.

B is running a route. He faces and moves the same way, at about 7.5 yards a second, heading toward the offense's right (**dir** well above 90 means he is drifting toward low y, which after standardising is the offense's right): H, who was in motion at the snap and is now on his way to the flat.

C is standing still, or close to it: his speed is 0.1 yards a second, well under the 1-yard-a-second threshold. His **o** is meaningful (he faces downfield, set to block), but his **dir** is just the direction of the chips' wobble and could be anything. That is the running back, who stepped up beside the left guard and is waiting for a blitzer who never came. Any feature built on **dir** should drop rows like C's: filter on **s** first.